

# Adaptive Mesh Refinement in Chapel: An Acid Test for High Productivity Programming

Jonathan Claridge  
Cray Inc. Tech Forum  
December 2, 2010

# Overview

## Goal

- Initially modest: Isolate “motifs” of adaptive mesh refinement, suitable for benchmarking

# Overview

## Goal

- Initially modest: Isolate “motifs” of adaptive mesh refinement, suitable for benchmarking

## Results

- Fully-functional, **dimension-independent** AMR framework in under 4 months, with no prior Chapel experience

# Overview

## Goal

- Initially modest: Isolate “motifs” of adaptive mesh refinement, suitable for benchmarking

## Results


- Fully-functional, **dimension-independent** AMR framework in under 4 months, with no prior Chapel experience
- Code is drastically shorter than existing libraries:

Language	Parallelism	SLOC <sup>1</sup>	Tokens	Relative size (tokens)
Chapel (any D)	Shared mem.	1988	13783	1
Fortran (2D+3D) <sup>2</sup>	Serial	16562	151992	11.03
2D		8297	71639	5.20
3D		8265	80353	5.83
C/C++ (any D) <sup>3</sup>	Dist. mem.	40200	261427	20.22

<sup>1</sup> source lines of code, <sup>2</sup> AMRClaw, <sup>3</sup> Chombo BoxTools+AMRTools


# Adaptive Mesh Refinement (AMR)

Provide enhanced resolution in regions where features are poorly resolved




# Adaptive Mesh Refinement (AMR)

Provide enhanced resolution in regions where features are poorly resolved




# Adaptive Mesh Refinement (AMR)

Provide enhanced resolution in regions where features are poorly resolved




# AMR Terminology









# AMR Terminology



# AMR Terminology



# AMR Terminology



# Spatial dimension

```
config param dimension = 2;
const dimensions = 1..dimension;
```

# Spatial dimension

May specify value  
with a command-  
line flag

```
config param dimension = 2;
```

```
const dimensions = 1..dimension;
```

# Spatial dimension

May specify value  
with a command-  
line flag

Value known at  
compile-time

**config param** dimension = 2;

**const** dimensions = 1..dimension;

# Spatial dimension

May specify value  
with a command-  
line flag

Value known at  
compile-time

Default value;  
type is elided

**config param** dimension = 2;

**const** dimensions = 1..dimension;



# Spatial dimension

May specify value  
with a command-  
line flag

Value known at  
compile-time

Default value;  
type is elided

**config param** dimension = 2;

**const** dimensions = 1..dimension;

Not modified after  
assignment;  
use **var** otherwise

# Spatial dimension

May specify value  
with a command-  
line flag

Value known at  
compile-time

Default value;  
type is elided

**config param** dimension = 2;

**const** dimensions = 1..dimension;

Not modified after  
assignment;  
use **var** otherwise

**Range:** Arithmetic sequence of integers


- Supports iteration:

**for** d **in** dimensions **do** ...

# Grids

Basic geometry described by tuples

```
const x_low, x_high: dimension*real;
```




# Grids

## Basic geometry described by tuples

```
const x_low, x_high: dimension*real;
```

```
const i_low, i_high: dimension*int;
```



# Grids

## Basic geometry described by tuples


```
const x_low, x_high: dimension*real;
```

```
const i_low, i_high: dimension*int;
```

```
const n_cells: dimension*int;
```

n\_cells(2)=5

n\_cells(1)=4



# Grids

## Basic geometry described by tuples

```
const x_low, x_high: dimension*real;
```


```
const i_low, i_high: dimension*int;
```

```
const n_cells: dimension*int;
```

```
const n_ghost_cells: dimension*int;
```

```
n_ghost_cells(2)=2
```

```
n_ghost_cells(1)=1
```



# Grids: Indexing

Conventional approach – number the cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

# Grids: Indexing

Conventional approach – number the cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```



# Grids: Indexing

Conventional approach – number the cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```

## Arithmetic domain: Multidimensional index space

- Supports storage:

```
var my_array: [cells] real;
```

- Supports (parallel) iteration:

```
for(all) cell in cells do ...
```

- The reason Chapel is so useful for AMR

# Grids: Indexing

Conventional approach – number the cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```

Problem with conventional indexing:

- How are interfaces and vertices indexed?

# Grids: Indexing


Conventional approach – number the cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```

Problem with conventional indexing:


- How are interfaces and vertices indexed?



Many objects will  
have the same indices


# Grids: Indexing

Modified approach – denser index space



# Grids: Indexing


Modified approach – denser index space



**const cells** = [1..7 **by** 2, 1..5 **by** 2];

# Grids: Indexing

Modified approach – denser index space




```
const cells          = [1..7 by 2, 1..5 by 2];
```

```
const x_interfaces = [0..8 by 2, 1..5 by 2];
```

# Grids: Indexing


Modified approach – denser index space



```
const cells          = [1..7 by 2, 1..5 by 2];
const x_interfaces = [0..8 by 2, 1..5 by 2];
const y_interfaces = [1..7 by 2, 0..6 by 2];
```

# Grids: Indexing

Modified approach – denser index space



```
const cells           = [1..7 by 2, 1..5 by 2];
const x_interfaces = [0..8 by 2, 1..5 by 2];
const y_interfaces = [1..7 by 2, 0..6 by 2];
const vertices      = [0..8 by 2, 0..6 by 2];
```



# Grids: Indexing

Modified approach – denser index space



## ***Strided domains***


- Array and iteration syntax are **unchanged**
- Chapel helps us describe the physical problem *much* more effectively

```

const cells =          [1..7 by 2, 1..5 by 2];
const x_interfaces = [0..8 by 2, 1..5 by 2];
const y_interfaces = [1..7 by 2, 0..6 by 2];
const vertices =      [0..8 by 2, 0..6 by 2];
  
```

# Levels


Essentially a union of grids



# Levels


Essentially a union of grids

```
var grids: domain(Grid);
```



# Levels

Essentially a union of grids




```
var grids: domain(Grid);
```

## Associative domain

- A list of indices of *any* type
- Array and iteration syntax remains **unchanged**

# Levels: Overlaps


Each grid has a layer of **ghost cells** to facilitate data transfer



# Levels: Overlaps

Each grid has a layer of **ghost cells** to facilitate data transfer

```
const extended_cells =  
    cells.expand(n_ghost_cells);
```




# Levels: Overlaps

Each grid has a layer of **ghost cells** to facilitate data transfer

```
const extended_cells =  
    cells.expand(n_ghost_cells);
```


Calculating the **overlaps**



# Levels: Overlaps

Each grid has a layer of **ghost cells** to facilitate data transfer

```
const extended_cells =  
    cells.expand(n_ghost_cells);
```



## Calculating the overlaps

```
var neighbors: domain(Grid);  
var overlap = extended_cells( sibling.cells, dimension, stridable=true);  
for sibling in parent_level.grids {  
    var overlap = extended_cells( sibling.cells );  
  
    if overlap.numIndices > 0 && sibling != this {  
        neighbors.add(sibling);  
        overlaps(sibling) = overlap;  
    }  
}
```


Declare associative domain to store neighbors; initializes to empty.



# Levels: Overlaps

Each grid has a layer of **ghost cells** to facilitate data transfer

```
const extended_cells =  
    cells.expand(n_ghost_cells);
```



## Calculating the overlaps

```
var neighbors: domain(Grid);  
var overlaps: [neighbors] domain(dimension, stridable=true);
```


```
for sibling in neighbors:  
    var overlap = ...;  
  
if overlap.num_cells > 0 {  
    neighbors.add(sibling);  
    overlaps(sibling) = overlap;  
}
```

An array of domains; stores one domain for each neighbor.  
New space allocated as neighbors grows.

# Levels: Overlaps

Each grid has a layer of **ghost cells** to facilitate data transfer

```
const extended_cells =  
    cells.expand(n_ghost_cells);
```



## Calculating the overlaps


```
var neighbors: domain(Grid);  
var overlaps: [neighbors] domain(dimension, stridable=true);  
  
for sibling in parent_level.grids {  
    var overlap = sibling.cells;  
    if overlap != this {  
        neighbors.add(sibling);  
        overlaps(sibling) = overlap;  
    }  
}
```

Loop over all grids on the same level, checking for neighbors.

# Levels: Overlaps

Each grid has a layer of **ghost cells** to facilitate data transfer

```
const extended_cells =  
    cells.expand(n_ghost_cells);
```



## Calculating the overlaps


```
var neighbors: domain(Grid);  
var overlaps: [neighbors] domain(dimension, stridable=true);  
  
for sibling in parent_level.grids {  
    var overlap = extended_cells( sibling.cells );  
  
    if overlap {  
        neighbors[sibling] = overlap;  
        overlaps(sibling) = overlap;  
    }  
}
```

Computes intersection of the domains  
extended\_cells and sibling.cells.

# Levels: Overlaps

Each grid has a layer of **ghost cells** to facilitate data transfer

```
const extended_cells =  
    cells.expand(n_ghost_cells);
```




## Calculating the overlaps

```
var neighbors: domain(Grid);  
var overlaps: [neighbors] domain(dimension, stridable=true);  
  
for sibling in parent_level.grids {  
    var overlap = extended_cells( sibling.cells );  
  
    if overlap.numIndices > 0 && sibling != this {  
        neighbors.add(sibling);  
        overlaps(sibling) = overlap;  
    }  
}
```

If overlap is nonempty, and sibling is distinct from this grid, then update stored data.

# Levels: Interpolation


Remaining ghost cells will receive data  
interpolated from the coarser level



# Levels: Interpolation

Remaining ghost cells will receive data  
**interpolated from the coarser level**


- Region is naturally described by a union of domains



# Levels: Interpolation

Remaining ghost cells will receive data  
**interpolated from the coarser level**


- Region is naturally described by a union of domains
- Region is naturally calculated by *subtraction* of domains



# Levels: Interpolation

Remaining ghost cells will receive data  
**interpolated from the coarser level**

- Region is naturally described by a union of domains
- Region is naturally calculated by *subtraction* of domains



## New object: **MultiDomain**


- Stores a collection of domains
- Supports subtraction of domains with a set-minus interpretation



# Levels: Interpolation

Remaining ghost cells will receive data  
**interpolated from the coarser level**

- Region is naturally described by a union of domains
- Region is naturally calculated by *subtraction* of domains




Computing the interpolation region

# Levels: Interpolation

Remaining ghost cells will receive data  
**interpolated from the coarser level**

- Region is naturally described by a union of domains
- Region is naturally calculated by *subtraction* of domains



## Computing the interpolation region

```
var interp_region =  
    new MultiDomain(dimension, stridable=true);
```

```
int  
int
```


Declare a new MultiDomain; initializes to empty.

```
for neighbor in neighbors do  
    interp_region.subtract( overlaps(neighbor) );
```

# Levels: Interpolation

Remaining ghost cells will receive data  
**interpolated from the coarser level**

- Region is naturally described by a union of domains
- Region is naturally calculated by *subtraction* of domains



## Computing the interpolation region

```
var interp_region =  
    new MultiDomain(dimension, stridable=true);  
interp_region.add( extended_cells );  
interp_region.subtract( cells );
```


Fill with **only** the ghost cells.

```
for neighbor in neighbors do  
    interp_region.subtract( overlaps(neighbor) );
```

# Levels: Interpolation

Remaining ghost cells will receive data  
**interpolated from the coarser level**

- Region is naturally described by a union of domains
- Region is naturally calculated by *subtraction* of domains



## Computing the interpolation region

```
var interp_region =
    new MultiDomain(dimension, stridable=true);
interp_region.add( extended_cells );
interp_region.subtract( cells );
```


```
for neighbor in neighbors do
    interp_region.subtract( overlaps(neighbor) );
```

Remove any regions of overlap with a neighboring grid.

# Levels: Interpolation

Remaining ghost cells will receive data interpolated from the coarser level

- Region is naturally described by a union of domains




- This code is dimension-independent
- Underlying code for MultiDomains is relatively simple because domains do most of the hard work

```
var interp_region =  
    new MultiDomain(dimension, stridable=true);  
interp_region.add( extended_cells );  
interp_region.subtract( cells );
```

```
for neighbor in neighbors do  
    interp_region.subtract( overlaps(neighbor) );
```


# Regridding (very briefly)

- Primarily the work of a **partitioning algorithm**: Given a set of flags, find rectangles that cover them (Berger & Rigoutsos, 1991)



# Regridding (very briefly)

- Primarily the work of a **partitioning algorithm**: Given a set of flags, find rectangles that cover them (Berger & Rigoutsos, 1991)




- More rectangles and unions of rectangles – naturally described by domains and MultiDomains



# Regridding (very briefly)

- Primarily the work of a **partitioning algorithm**: Given a set of flags, find rectangles that cover them (Berger & Rigoutsos, 1991)



- More rectangles and unions of rectangles – naturally described by domains and MultiDomains
- Dimension-independent code is  $\approx 200$  lines



# Parallelism

## Shared-memory parallelism

- Genuinely trivial:

```
forall cell in grid.cells do ...
```

# Parallelism

## Shared-memory parallelism

- Genuinely trivial:

```
forall cell in grid.cells do ...
```

## Distributed-memory parallelism

- Not implemented yet, but not difficult once grids have been mapped to processors

```
forall grid in level.grids {
  forall cell in grid.cells do ...
}
```

# Performance

- Chapel performance in general is an open issue

# Performance

- Chapel performance in general is an open issue
- Will at least test:
  - What fraction of time is spent number crunching?
  - How good is scaling on a multicore machine?

# Summary

- The challenges of AMR are mostly gymnastics with rectangles. Chapel makes this immensely easier.

# Summary

- The challenges of AMR are mostly gymnastics with rectangles. Chapel makes this immensely easier.
- Dimension-independence and parallelism are so simple that you can almost forget about them.

# Summary

- The challenges of AMR are mostly gymnastics with rectangles. Chapel makes this immensely easier.
- Dimension-independence and parallelism are so simple that you can almost forget about them.
- Massive reductions in code size:

Language	Parallelism	SLOC <sup>1</sup>	Tokens	Relative size (tokens)
Chapel (any D)	Shared mem.	1988	13783	1
Fortran (2D+3D) <sup>2</sup>	Serial	16562	151992	11.03
2D		8297	71639	5.20
3D		8265	80353	5.83
C/C++ (any D) <sup>3</sup>	Dist. mem.	40200	261427	20.22

<sup>1</sup> source lines of code, <sup>2</sup> AMRClaw, <sup>3</sup> Chombo BoxTools+AMRTools