

HPCC STREAM and RA in Chapel

Performance and Potential

Steve Deitz

Cray Inc.

What is Chapel?

- A new parallel language
 - Under development at Cray Inc.
 - Supported through the DARPA HPCS program
- Goals
 - **Improve programmer productivity**
 - Improve the programmability of parallel computers
 - Match or improve performance of MPI/UPC/CAF
 - Provide better portability than MPI/UPC/CAF
 - Improve robustness of parallel codes
 - Support multi-core and multi-node systems


Outline

- What is Chapel?
- Chapel's Parallel Programming Model
- HPCC STREAM Triad in Chapel
- HPCC RA in Chapel
- Summary and Future Work

Fragmented vs. Global-View: Definitions


- Programming model
The mental model of a programmer
- Fragmented models
Programmers take point-of-view of a single processor/thread
- SPMD models (Single Program, Multiple Data)
Fragmented models with multiple copies of one program
- Global-view models
Programmers write code to describe computation as a whole

3-Point Stencil Example (n=6)



3-Point Stencil Example


Global-View vs. Fragmented Computation




3-Point Stencil Example: Code

Global-View vs. Fragmented Code

Global-View



Fragmented



Assumes p divides n

```
def main() {
  var n = 1000;
  var me = commID(), p = commProcs(),
        myN = n/p, myLo = 1, myHi = myN;
  var A, B: [0..myN+1] real;

  if me < p {
    send(me+1, A(myN));
    recv(me+1, A(myN+1));
  } else myHi = myN-1;
  if me > 1 {
    send(me-1, A(1));
    recv(me-1, A(0));
  } else myLo = 2;
  for i in myLo..myHi do
    B(i) = (A(i-1)+A(i+1))/2;
  }
```

```
def rprj3(S, R) {
  const Stencil = [-1..1, -1..1, -1..1],
    W: [0..3] real = (0.5, 0.25, 0.125, 0.0625),
    W3D = [(i,j,k) in Stencil] W((i!=0)+(j!=0)+(k!=0));
```

[illegible]

Outline

- What is Chapel?
- Chapel's Parallel Programming Model
- HPCC STREAM Triad in Chapel
- HPCC RA in Chapel
- Summary and Future Work

Introduction to STREAM Triad

Given: m -element vectors A, B, C


Compute: `forall i in 1..m do`
 $A(i) = B(i) + \alpha * C(i);$



Distributed Parallelization of STREAM Triad

Given: m -element vectors A, B, C

Compute: `forall i in 1..m do`
 $A(i) = B(i) + \alpha * C(i);$




Further Parallelization of STREAM Triad

Given: m -element vectors A, B, C

Compute: `forall i in 1..m do`

$$A(i) = B(i) + \alpha * C(i);$$



STREAM Triad in Chapel: Single Locale

Given: m -element vectors A, B, C

Compute: `forall i in 1..m do`
 $A(i) = B(i) + \alpha * C(i);$

```

config const m: int(64) = ...;
const alpha: real = 3.0;
const ProblemSpace: domain(1,int(64)) = [1..m];
var A, B, C: [ProblemSpace] real;

forall i in ProblemSpace do
  A(i) = B(i) + alpha * C(i);
  
```

STREAM Triad in Chapel: Single Locale

Given: m -element vectors A, B, C

Compute: `forall i in 1..m do`
`A(i) = B(i) + α * C(i);`

```
config const m: int(64) = ...;
const alpha: real = 3.0;
const ProblemSpace: domain(1,int(64)) = [1..m];
var A, B, C: [ProblemSpace] real;
```

`A = B + alpha * C;`

More concise variation
using whole array operations

STREAM Triad in Chapel: Single Locale

Given: m -element vectors A, B, C

Compute: `forall i in 1..m do`
 $A(i) = B(i) + \alpha * C(i);$

```
config const m: int(64) = ...;
const alpha: real = 3.0;
const ProblemSpace: domain(1,int(64)) = [1..m];
var A, B, C: [ProblemSpace] real;
```

```
forall (a,b,c) in (A,B,C) do
  a = b + alpha * c;
```

Variation that iterates
directly over the arrays

STREAM Triad in Chapel: Multi-Locale

Given: m -element vectors A, B, C

Compute: `forall i in 1..m do`

`A(i) = B(i) + α * C(i);`

```
config const m: int(64) = ..., tpl = ...;
```

```
const alpha: real = 3.0;
```

```
const BlockDist = new Block(1, int(64), [1..m], tpl);
```

```
const ProblemSpace: domain(1, int(64))
```

```
distributed BlockDist = [1..m];
```

```
var A, B, C: [ProblemSpace] real;
```


```
forall (a,b,c) in (A,B,C) do
```

```
  a = b + alpha * c;
```



What is a Distribution?

A “recipe” for distributed arrays that...

Instructs the compiler how to Map the global view...



...to a fragmented, per-processor implementation



STREAM Triad in Chapel: Multi-Locale

Given: m -element vectors A, B, C

Compute: `forall i in 1..m do`
 $A(i) = B(i) + \alpha * C(i);$


```

config const m: int(64) = ..., tpl = ...;
const alpha: real = 3.0;
const BlockDist = new Block(1,int(64), [1..m], tpl);
const ProblemSpace: domain(1, int(64))
                        distributed BlockDist = [1..m];
var A, B, C: [ProblemSpace] real;


forall (a,b,c) in (A,B,C) do
  a = b + alpha * c;
  
```

HPCC STREAM Performance

Performance of STREAM in Chapel



HPCC STREAM Efficiency



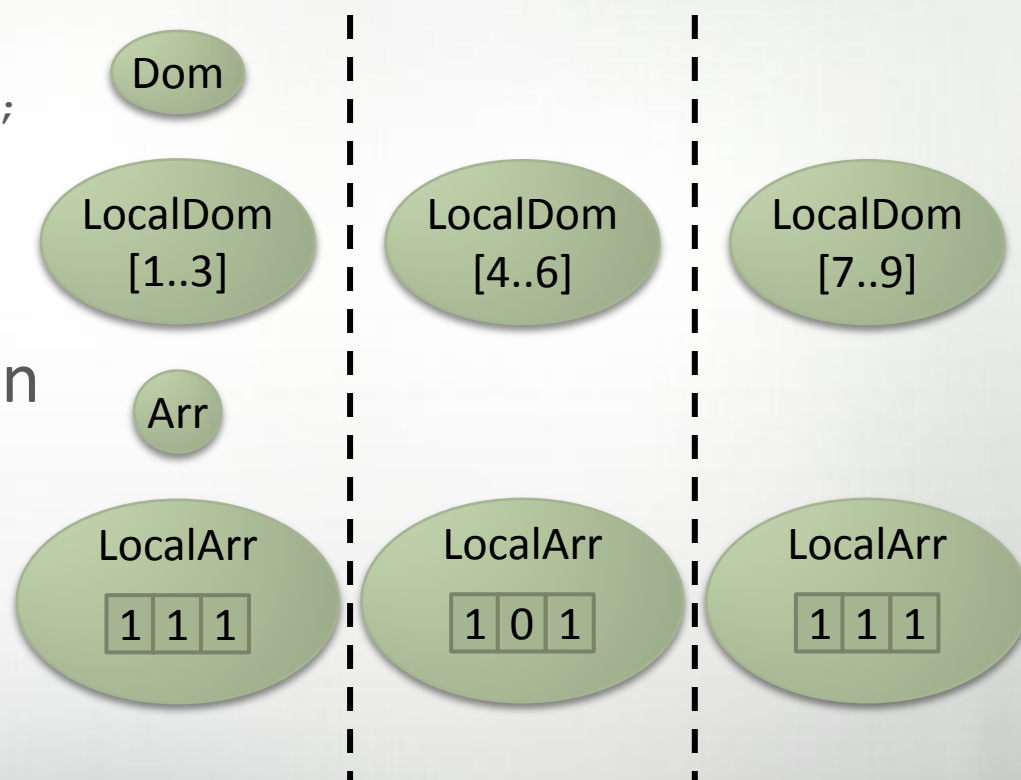
Optimization: Privatization

Simple example

```
var Dist: Block(1,int(64));
var Dom: domain(1,int(64))
    distributed Dist;
var Arr: [Dom] int;
```

Reference to local data
requires communication

```
on Locales(1) {
  Arr(5) = 0;
}
```



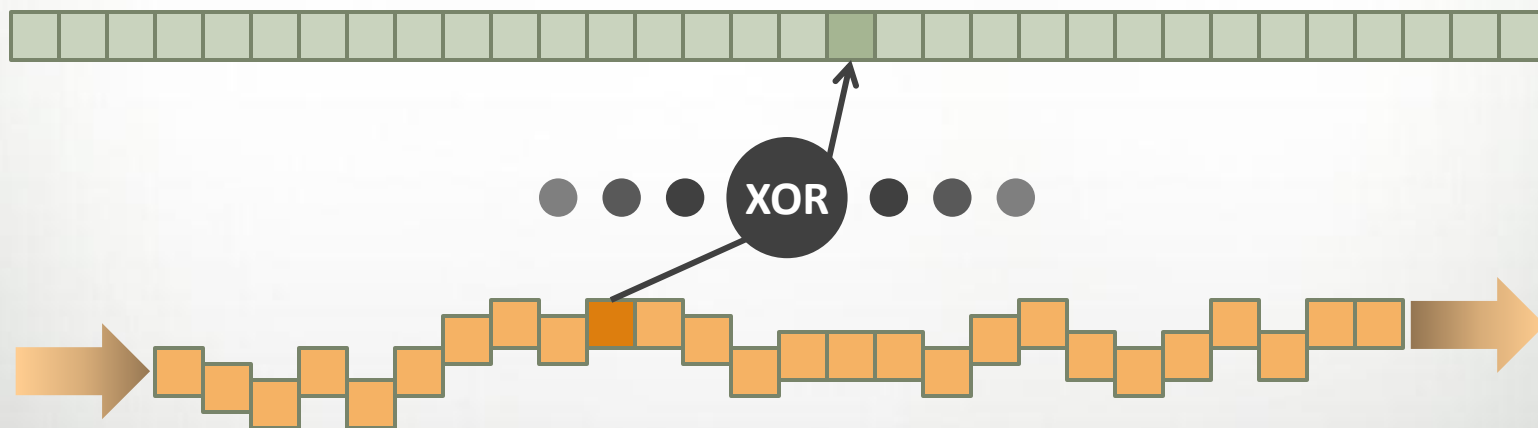
Outline

- What is Chapel?
- Chapel's Parallel Programming Model
- HPCC STREAM Triad in Chapel
- HPCC RA in Chapel
- Summary and Future Work

Introduction to RA

Given: m -element table T (where $m = 2^n$)


Compute: `forall r in RandomUpdates do`
`$T(r \ \& \ (m-1)) \hat{=} r;$`



Distributed Parallelization of RA

Given: m -element table T (where $m = 2^n$)

Compute: `forall r in RandomUpdates do`
`$T(r \ \& \ (m-1)) \hat{=} r;$`



RA in Chapel: Single Locale

Given: m -element table T (where $m = 2^n$)

Compute: **forall** r **in** **RandomUpdates** **do**
 $T(r \ \& \ (m-1)) \ \wedge= \ r;$

```
config const m = ..., N_U = ...;
const TableSpace: domain(1,uint(64)) = [0.. $m-1$ ],
    Updates: domain(1,uint(64)) = [0.. $N_U-1$ ];
var T: [TableSpace] uint(64);

forall (i,r) in (Updates,RAStream()) do
    T(r & (m-1)) ^= r;
```

RA in Chapel: Multi-Locale

Given: m -element table T (where $m = 2^n$)

Compute: **forall** r **in** **RandomUpdates** **do**
 $T(r \ \& \ (m-1)) \ \wedge= \ r;$

```

config const m = ..., N_U = ..., tpl = ...;
const TableDist = new Block(1,uint(64),[0.. $m-1$ ],tpl),
    UpdateDist = new Block(1,uint(64),[0.. $N\_U-1$ ],tpl),
    TableSpace: domain(1,uint(64))
                distributed TableDist = [0.. $m-1$ ],
    Updates: domain(1,uint(64))
            distributed UpdateDist = [0.. $N\_U-1$ ];
var T: [TableSpace] uint(64);

forall (i,r) in (Updates,RAStream()) do
    on T( $r \ \& \ (m-1)$ ) do
        T( $r \ \& \ (m-1)$ )  $\wedge= r;$ 
  
```

RA in Chapel: Multi-Locale

Given: m -element table T (where $m = 2^n$)

Compute: `forall r in RandomUpdates do`
`T(r & (m-1)) ^= r;`

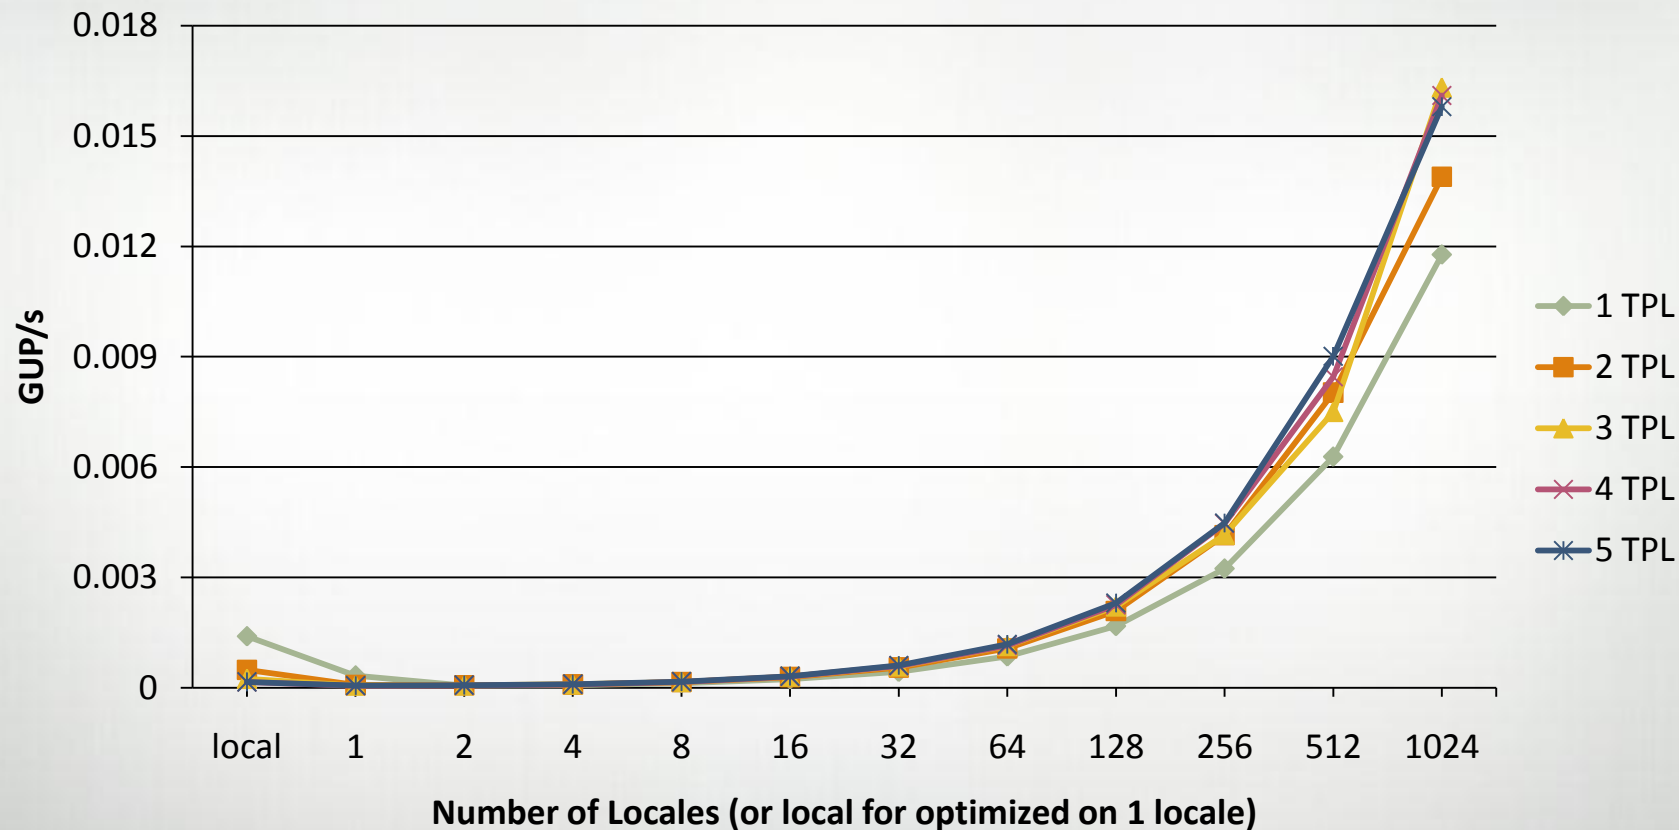
```
config const m = ..., N_U = ..., tpl = ...;
const TableDist = new Block(1, uint(64), [0..m-1], tpl),
      UpdateDist = new Block(1, uint(64), [0..N_U-1], tpl),
      TableSpace: domain(1, uint(64))
                  distributed TableDist = [0..m-1],
      Updates: domain(1, uint(64))
                  distributed UpdateDist = [0..N_U-1];
var T: [TableSpace] uint(64);

forall (i, r) in (Updates, RASStream()) do
  on T.domain.dist.ind2loc(r & (m-1)) do
    T(r & (m-1)) ^= r;
```


Call ind2loc method directly

HPCC RA Performance

Performance of RA in Chapel

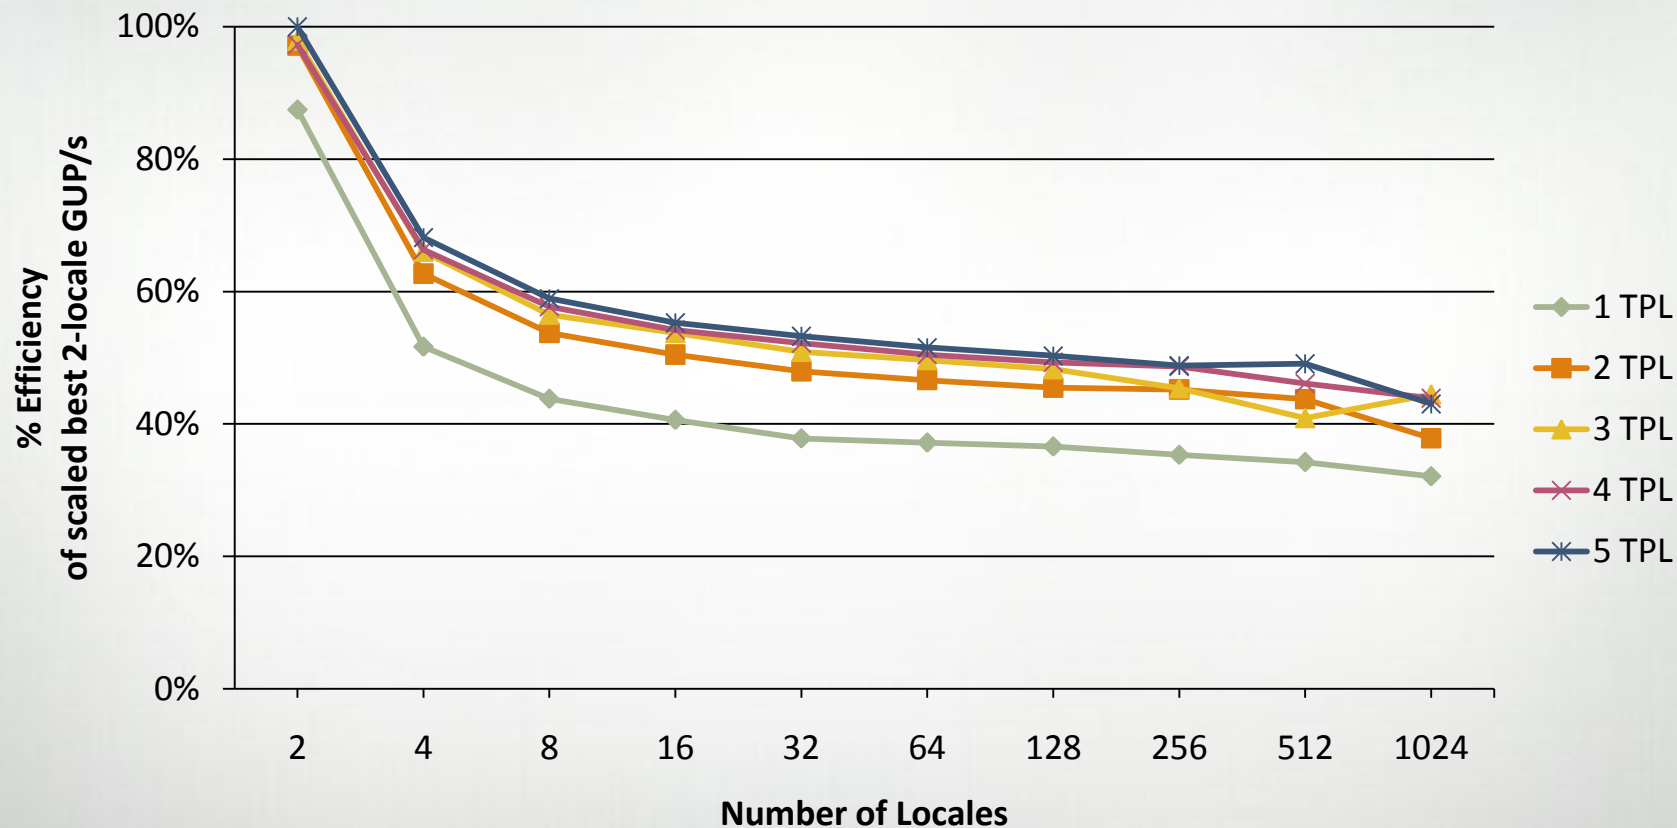


HPCC RA Efficiency I



HPCC RA Efficiency II


Efficiency of RA in Chapel



Optimization: Remote Value Forwarding

Simple example

```
var Arr: [Dom] int;
var r: int;
on Locales(1) {
  Arr(r) ^= r;
}
```



Outline

- What is Chapel?
- Chapel's Parallel Programming Model
- HPCC STREAM Triad in Chapel
- HPCC RA in Chapel
- Summary and Future Work

Summary

The global-view programming model is easy to use.

- Shorter, more concise code
- Separation of concerns (partitioning)
- Easy to change data distributions

Distributions implement the global-view model.

- Flexible mechanism for experimentation
- Implementation of distributions is in Chapel

Future Work

- Optimizations
 - Within the compiler
 - Within the runtime
 - Within the distributions
- Complete implementation of Block distribution
- Implement new distributions
 - Cyclic, BlockCyclic, RecursiveBisection
- Experiment with variations of STREAM and RA

Questions?