

# Hierarchical Locales: Exposing the Node Architecture in Chapel

Sung-Eun Choi, Cray Inc.

2<sup>nd</sup> KIISE-KOCSEA SIG HPC Workshop

November 19, 2013




# STREAM Triad: a trivial parallel computation

**Given:**  $m$ -element vectors  $A, B, C$

**Compute:**  $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

**In pictures:**



# STREAM Triad: MPI



## MPI

```
#include <hpcc.h>
```


```
static int VectorSize;  
static double *a, *b, *c;
```

```
int HPCC_StarStream(HPCC_Params *params) {  
    int myRank, commSize;  
    int rv, errCount;  
    MPI_Comm comm = MPI_COMM_WORLD;  
  
    MPI_Comm_size( comm, &commSize );  
    MPI_Comm_rank( comm, &myRank );  
  
    rv = HPCC_Stream( params, 0 == myRank );  
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,  
        0, comm );  
  
    return errCount;  
}
```

```
int HPCC_Stream(HPCC_Params *params, int doIO) {  
    register int j;  
    double scalar;
```

```
    VectorSize = HPCC_LocalVectorSize( params, 3,  
        sizeof(double), 0 );
```

```
    a = HPCC_XMALLOC( double, VectorSize );  
    b = HPCC_XMALLOC( double, VectorSize );  
    c = HPCC_XMALLOC( double, VectorSize );
```



```
    if (!a || !b || !c) {  
        if (c) HPCC_free(c);  
        if (b) HPCC_free(b);  
        if (a) HPCC_free(a);  
        if (doIO) {  
            fprintf( outFile, "Failed to allocate memory  
                (%d).\n", VectorSize );  
            fclose( outFile );  
        }  
        return 1;  
    }
```


```
    for (j=0; j<VectorSize; j++) {  
        b[j] = 2.0;  
        c[j] = 0.0;  
    }
```

```
    scalar = 3.0;
```

```
    for (j=0; j<VectorSize; j++)  
        a[j] = b[j]+scalar*c[j];
```

```
    HPCC_free(c);  
    HPCC_free(b);  
    HPCC_free(a);
```

# STREAM Triad: MPI+OpenMP



## MPI + OpenMP

```
#include <hpcc.h>
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM,
        0, comm );

    return errCount;
}

int HPCC_Stream(HPCC_Params *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
        sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
```


```
    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory
(%d).\n", VectorSize );
            fclose( outFile );
        }
        return 1;
    }
```

```
#ifdef _OPENMP
#pragma omp parallel for
#endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 0.0;
    }

    scalar = 3.0;
```

```
#ifdef _OPENMP
#pragma omp parallel for
#endif
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];

    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);
```



# STREAM Triad: MPI+OpenMP vs. CUDA

## MPI + OpenMP

```
#ifndef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );

    return errCount;
}

int HPCC_Stream(HPCC_Params *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );

    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
            fclose( outFile );
        }
        return 1;
    }


    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 0.0;
    }

    scalar = 3.0;

    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];

    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);

    return 0;
}
```



## CUDA

```
#define N 2000000

int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**) &d_a, sizeof(float)*N);
    cudaMalloc((void**) &d_b, sizeof(float)*N);
    cudaMalloc((void**) &d_c, sizeof(float)*N);

    dim3 dimBlock(128);
    dim3 dimGrid(N/dimBlock.x);
    if( N % dimBlock.x != 0 ) dimGrid


    set_array<<<dimGrid,dimBlock>>>(d_b, .5f, N);
    set_array<<<dimGrid,dimBlock>>>(d_c, .5f, N);

    scalar=3.0f;
    STREAM_Triad<<<dimGrid,dimBlock>>>(d_b, d_c, d_a, scalar, N);
    cudaThreadSynchronize();

    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

    __global__ void set_array(float *a, float value, int len) {
        int idx = threadIdx.x + blockIdx.x * blockDim.x;
        if (idx < len) a[idx] = value;
    }

    __global__ void STREAM_Triad( float *a, float *b, float *c,
                                float scalar, int len) {
        int idx = threadIdx.x + blockIdx.x * blockDim.x;
        if (idx < len) c[idx] = a[idx]+scalar*b[idx];
    }
}
```



# STREAM Triad: Chapel

## MPI + OpenMP

```
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params,
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT,
        0, comm );
    return errCount;
}

int HPCC_Stream(HPCC_Params *params,
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );

    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {

```

```
config const m = 1000,
    alpha = 3.0;


const ProblemSpace = {1..m} dmapped ...;

var A, B, C: [ProblemSpace] real;

B = 2.0;
C = 3.0;

A = B + alpha * C;
```

the special sauce



Philosophy: Good language design can tease details of locality and parallelism away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.

# Outline

- ✓ Motivation
- Chapel Background
  - Exposing the node architecture: locale models
  - Project Status and Next Steps

# What is Chapel?

- **An emerging parallel programming language**
  - Design and development led by Cray Inc.
    - in collaboration with academia, labs, industry
  - Initiated under the DARPA HPCS program
- **Overall goal: Improve programmer productivity**
  - Improve the **programmability** of parallel computers
  - Match or beat the **performance** of current programming models
  - Support better **portability** than current programming models
  - Improve the **robustness** of parallel codes
- **A work-in-progress**



# Chapel's Implementation

- Being developed as open source at SourceForge
- Licensed as BSD software
- **Target Architectures:**
  - multicore desktops and laptops
  - Cray architectures
  - commodity clusters
  - systems from other vendors
  - *in-progress*: CPU+accelerator hybrids, ...

# Chapel Design Goal

Write code like this...

```
config const m = 1000,
              alpha = 3.0;

const ProblemSpace = {1..m} dmapped ...;

var A, B, C: [ProblemSpace] real;

B = 2.0;
C = 3.0;

A = B + alpha * C;
```

Have it perform like this...

```
#include <stdio.h>
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );

    return errCount;
}

int HPCC_Stream(HPCC_Params *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );

    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
            fclose( outFile );
        }
    }

    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 0.0;
    }

    scalar = 3.0;

    #ifdef _OPENMP
    #pragma omp parallel for
    #endif
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];
    ...
}
```

...on any architecture available



# How can this be achieved?

## Don't expose low-level mechanisms

*Chapel provides syntax to express locality and parallelism, which are portable to any architecture*

## Don't expect the compiler to do everything

*Chapel's locality and parallelism mechanisms are implemented in Chapel modules not our compiler*

## Don't keep the implementation too close

*Chapel is an open-source project with contributors from around the world*

# Outline

- ✓ Motivation
- ✓ Chapel Background
- Exposing the node architecture: locale models
- Project Status and Next Steps

# Exposing the node architecture: Locale models



- **Locale models**
  - The locale type
  - Locality/affinity control
  - Task parallelism
- **Example: The NUMA locale model**
- **STREAM Triad revisited**
  - Domain maps
  - Leader-follower iterators

# The Locale Model

- Every Chapel program employs a locale model to describe the system architecture
- **Locale models define:**
  - Abstract node architecture (the *locale* type)
  - Memory allocation
  - Task scheduling policies
  - Locality/affinity control
- **Locale models are written in Chapel**
  - Class and procedural interface for the compiler
  - Utilizes runtime interfaces (e.g., communication, tasking/threading)
  - Programmer role: Architecture specialist

# The locale type

- **The *locale* type**
  - an abstract unit of architecture
    - Can have memory, processing units, etc.
  - can be nested (hierarchical)
  - typically a compute node (multicore processor or SMP) or your laptop
- **Chapel programs run on one or more locales**



# Defining Locales

- Specify # of locales when running Chapel programs

```
% a.out --numLocales=8
```

```
% a.out -nl 8
```

- Chapel provides built-in locale variables

```
config const numLocales: int = ...;  
const Locales: [0..#numLocales] locale = ...;
```

*Locales*

L0	L1	L2	L3	L4	L5	L6	L7
----	----	----	----	----	----	----	----

- User's `main()` begins executing on locale 0



# Locale Operations

- Locale methods support queries about the target system:

```
proc locale.numCores { ... }
proc locale.id { ... }
proc locale.name { ... }
```

- *On* statements support placement of computations:

```
writeln("on locale 0");

on Locales[1] do
  writeln("now on locale 1");

writeln("on locale 0 again");
```



# Locale Operations

- Locales can be nested:

```
proc locale.getChild(...) { ... }  
proc locale.getChildCount { ... }
```

- Children (sublocales) are locales
- *On* statements support placement of computations on sublocales:

```
writeln("on locale 0");  
  
on here.getChild(1) do  
    writeln("now on locale 0, sublocale 1");  
  
writeln("on locale 0 again");
```

# On statements: Rewriting

*Conceptually*, the Chapel compiler translates:

```
on here.getChild(1) do
  writeln("now on locale 0, sublocale 1");
```

into:

```
proc on_fn() {
  writeln("now on locale 0, sublocale 1");
}

chpl_executeOn(here.getChild(1), on_fn, ...);
```

**chpl\_executeOn() is defined by the Locale Model**

# Task parallelism

- **Tasks are the basic unit of parallelism**
- **Tasks can be created by the following Chapel statements:**
  - begin
  - cobegin
  - coforall

# Task parallelism: begin statements

```
// create a fire-and-forget task for a statement  
begin writeln("hello world");  
writeln("good bye");
```

## Possible outputs:

```
hello world  
good bye
```

```
good bye  
hello world
```

# Task parallelism: cobegin statements

```
// create a task per child statement  
cobegin {  
    producer(1);  
    producer(2);  
    consumer(1);  
} // implicit join of the three tasks here
```



# Task parallelism: coforall loops

```
// create a task per iteration  
coforall t in 0..#numTasks {  
    writeln("Hello from task ", t, " of ", numTasks);  
} // implicit join of the numTasks tasks here  
  
writeln("All tasks done");
```

## Sample output:

```
Hello from task 2 of 4  
Hello from task 0 of 4  
Hello from task 3 of 4  
Hello from task 1 of 4  
All tasks done
```

# Begin statement: Rewriting

*Conceptually*, the Chapel compiler translates:

```
begin writeln("hello world");
```

into:

```
proc begin_fn() {
    writeln("hello world");
}


chpl_taskListAddBegin(begin_fn, ...);
```

**chpl\_taskListAddBegin()** is defined by the Locale Model



# Example: The NUMA locale model


- The first interesting Locale Model prototype
- $n$  NUMA domains per locale (sublocale)
  - $m$  cores per NUMA domain
- Memory divided between NUMA domains
  - shared by both NUMA domains



# Example: The NUMA locale model


```
class LocaleModel: AbstractLocaleModel {
  var numSublocales: int;
  var childLocales: [0..#numSublocales] NumaDomain;
  ...
  proc getChild(i) return childLocales[i];
  ...
}
```

```
class NumaDomain: AbstractLocaleModel {
  ...
  proc getChild(i) return nil;
  ...
}
```




## Example: `chpl_executeOn()`

```
// no off-node case, locale ids simplified
proc chpl_executeOn(loc, fn, args) {
  if (loc == here) {
    // run directly on this numa domain
    chpl_fhtable_call(fn, args);
  } else {
    // move to a different numa domain
    var orig = here.sublocid;
    chpl_task_setSubloc(loc.sublocid);
    chpl_fhtable_call(fn, args);
    chpl_task_setSubloc(orig);
  }
}
```




## Example: `chpl_taskListAddBegin()`

```
// Variant 1:
// Start the task on the specified sublocale
// Let the tasking layer decide what to do with "any"
proc chpl_taskListAddBegin(loc, fn, args) {
    chpl_task_addToTaskList(fn, args, loc.sublocid, ...);
}
```



# Example: chpl\_taskListAddBegin()


```
// Variant 1:
// Variant 2:
// Start the task on the specified sublocale
// Schedule "any" in a round robin fashion
proc chpl_taskListAddBegin(loc, fn, args) {
  const sublocid
    = if loc.sublocid != subloc_any then
      loc.sublocid; // use specified subloc
    else
      loc.nextSubloc.fetchAdd(1) % loc.numSublocales;
      // round robin using atomic count
  chpl_task_addToTaskList(fn, args, sublocid, ...);
}
```



# Example: chpl\_taskListAddBegin()

```
// Variant 1:
// Variant 2:
// Variant 3:
// Start the task on the specified sublocale
// Schedule "any" based on load
proc chpl_taskListAddBegin(loc, fn, args) {
  const sublocid
    = if loc.sublocid != subloc_any then
      loc.sublocid; // use specified subloc
    else
      getBestSubloc(loc);
  chpl_task_addToTaskList(fn, args, sublocid, ...);
}

proc getBestSubloc(loc) {
  const (, sublocid)
    = minloc reduce (loc.numTasks(), 0..#numSubLocs);
  return sublocid;
}
```






# Back to STREAM Triad

```
config const m = 1000,  
            alpha = 3.0;  
  
const ProblemSpace = {1..m} dmapped ...;  
  
var A, B, C: [ProblemSpace] real;  
  
B = 2.0;  
C = 3.0;  
  
A = B + alpha * C;
```


the special  
sauce

# The Special Sauce: Domain maps

Domain maps are “recipes” that instruct the compiler how to map the global view of a computation...



...to the target locales' memory and processors:





# Domain maps in a nutshell

- **Domain maps specify**
  - data allocation and layout
  - parallel iteration ←
  - low level implementation of array and domain operations
    - e.g., slicing, reallocating, reshaping
- **Domain maps are written in Chapel**
  - Class and procedural interface for the compiler
  - Utilizes lower level concepts like coforall and on statements
  - Programmer role: HPC specialist

# Back to STREAM Triad

```
config const m = 1000,
              alpha = 3.0;

const ProblemSpace = {1..m} dmapped ...;

var A, B, C: [ProblemSpace] real;

B = 2.0;
C = 3.0;

forall (a,b,c) in zip(A,B,C) do
    a = b + alpha * c;
```

- How many tasks?
- Where are they executed?
- How are the iterations assigned to tasks?

**A:** Chapel's *leader-follower* iterators are designed to give programmers full control over such decisions

# Leader-Follower Iterators: Definition

- Chapel defines all forall loops in terms of *leader-follower iterators*:
  - *leader iterators*: create parallelism, assign iterations to tasks
  - *follower iterators*: serially execute work generated by leader

- Given...

```
forall (a,b,c) in zip(A,B,C) do
    a = b + alpha * c;
```

...*A* is defined to be the *leader*

...*A*, *B*, and *C* are all defined to be *followers*

# Leader-follower iterators: Rewriting

Conceptually, the Chapel compiler translates:

```
forall (a,b,c) in zip(A,B,C) do
  a = b + alpha * c;
```




into:

```
coforall subloc in sublocales do on subloc do
  coforall tid in here.numCores {
    for (a,b,c) in zip(A,B,C) {
      a = b + alpha * c;
    }
  }
```

Leader

Followers



# Outline

- ✓ Motivation
- ✓ Chapel Background
- ✓ Exposing the node architecture: locale models
- Project Status and Next Steps

# Implementation Status -- Version 1.8.0 (Oct 2013)

## Overall Status:

- Most features work at a functional level
  - some features need to be improved or re-implemented (e.g., OOP)
- Many performance optimizations remain
  - particularly for distributed memory (multi-locale) execution

# Next Steps

- **Evolve from Prototype- to Production-grade**
  - Add/Improve missing features
  - Performance optimizations
- **Target more complex compute node types**
  - e.g., CPU+GPU, Intel Phi, ...
- **Continue to grow the user and developer communities**
  - Work toward transitioning Chapel from Cray-controlled to community-governed

# Chapel Design Goal Revisited

Write code like this...

```
config const m = 1000,
              alpha = 3.0;

const ProblemSpace = {1..m} dmapped ...;

var A, B, C: [ProblemSpace] real;

B = 2.0;
C = 3.0;

A = B + alpha * C;
```

Have it perform like this...

```
#include <stdio.h>
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );

    return errCount;
}

int HPCC_Stream(HPCC_Params *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );

    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
            fclose( outFile );
        }
    }
}
```

...on any architecture available

```
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 0.0;
}

scalar = 3.0;

#ifdef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];
...
```



# Chapel Design Goal Revisted

## ✓ High-level programmability

- Applications programmers expression parallelism and locality
- Chapel's locale models and domain maps can insulate programmers from details of the architecture

## ~ Performance

- Designed to enable good performance, results yet to be seen
- In cases where we've optimized, we've seen very promising results, e.g.,
  - Dynamic iterators that match the performance of OpenMP [Chamberlain, et al., PGAS 2011]
  - Communication aggregation implemented via the domain map interface [Sanz, et al., SBAC-PAD 2012]



# Chapel at SC13

- **Emerging Technologies Booth (all week)**
  - Booth #3547: staffed by Chapel team members; poster and handouts
- **Talk (Tues @ 3:20):** *Hierarchical Locales: Exposing the Node Architecture in Chapel*
  - KISTI booth (#3713): Sung-Eun Choi (Cray Inc.)
- **Poster (Tues @ 5:15):** *Towards Co-Evolution of Auto-Tuning and Parallel Languages*
  - Posters Session: Ray Chen (University of Maryland)
- **Chapel Lightning Talks BoF (Wed @ 12:15)**
  - 5-minute talks on education, MPI-3, Big Data, Autotuning, Futures, MiniMD
- **Talk (Wed @ 4:30):** *Chapel, an Emerging Parallel Language*
  - HPC Impact Theatre (booth #3947): Brad Chamberlain (Cray Inc.)
- **Happy Hour (Wed @ 5pm):** *4<sup>th</sup> annual Chapel Users Group (CHUG) Happy Hour*
  - Pi Bar (just across the street at 1400 Welton St): open to public, dutch treat
- **HPC Education (Thus @ 1:30pm):** *High-Level Parallel Programming Using Chapel*
  - David Bunde (Knox College) and Kyle Burke (Colby College)

