

Chapel

An Emerging Parallel Programming Language

If you wouldn't recommend your HPC programming model to a non-HPC friend, why aren't you demanding something better?

HPC has no languages that are as modern, productive, and enjoyable as Python / Matlab / Java / (your favorite language here). This is due less to technical challenges than to lack of community will, resources, effort, and patience.

Let's change that!

Chapel Characteristics

- General Parallelism
 - Any parallel algorithm
 - Any parallel hardware
- Separate Parallelism from Locality from Mechanism
 - What should run concurrently?
 - Where should tasks / data be placed?
 - Leave "how should it run?" to lower levels of software
- Designed for Performance
 - Features designed to compete with C/Fortran + MPI
- Modern Productivity Features
 - Type inference, iterators, rich array types, ranges, tuples, generic programming, optional OOP, ...

Chapel's Implementation


- Open Source
 - BSD License
 - Hosted and developed at SourceForge
- Scalable, Portable Design
 - From laptops to clusters, supercomputers, & the cloud
 - Requirements: C/C++, POSIX tools and threads
- Status
 - Version 1.8 released October 17, 2013
 - Core features are functional, ready to be used
 - Additional performance optimizations needed
- Community Effort
 - v1.8: 19 contributors from 8 organizations/5 countries

Multiresolution Design

Goal: Support multiple tiers of features

- higher levels for programmability, productivity
- lower levels for greater degrees of control

Chapel language concepts




- build the higher levels in terms of the lower
- permit the user to intermix layers arbitrarily

Domain Maps

Role: map global-view domains/arrays to locales

```
const ElemSpace = {0..#numElems}, local and distributed domains  
NodeSpace = {0..#numNodes}, (first-class index sets)  
Elems = ElemSpace dmapped Block(ElemSpace),  
Nodes = NodeSpace dmapped Block(NodeSpace);
```



Declarations from the Chapel version of LULESH

Data Parallelism

Role: drive parallelism via domains/arrays

```
// find the total velocity across atoms  
vtot = + reduce forall (bin, c) in zip(Bins, Count) do  
+ reduce forall a in bin[1..c] do a.v;
```

```
// compute the average velocity  
vtot /= numAtoms;
```

```
// adjust atom velocities using the average  
forall (bin, c) in zip(Bins, Count) do  
for a in bin[1..c] do  
a.v -= vtot;
```

zippered iteration is used to traverse multiple collections simultaneously

Computations from the Chapel version of MiniMD

Task Parallelism

Role: express parallelism via concurrent tasks

```
var buff$: [0..#buffsize] sync real;
```

```
cobegin {  
  producer();  
  consumer();  
}
```

```
proc producer() {  
  var i = 0;  
  for ... {  
    i = (i+1) % buffsize;  
    buff[i] = ...;  
  }
```

```
proc consumer() {  
  var i = 0;  
  while ... {  
    i = (i+1) % buffsize;  
    ...buff[i]...;  
  }
```

writes to sync variables block until empty, leave full

reads block until full, leave empty

Bounded buffer producer-consumer in Chapel

Base Language

Role: support productive serial programming

```
iter tiledRMO(D, tileSize) {  
  const tile = {0..#tilesize, 0..#tilesize};  
  for base in D by tileSize do  
    for ij in D[tile + base] do  
      yield ij;  
    }
```

the base language includes a rich algebra for operating on domains

yield a value back to the caller; then continue executing the iterator

```
for ij in tiledRMO({1..m, 1..n}, 2) do  
  write(ij);
```

This loop prints:
(1,1) (1,2) (2,1) (2,2) (1,3) (1,4) (2,3) (2,4) ...
(3,1) (3,2) (4,1) (4,2) (3,3) (3,4) (4,3) (4,4) ...

Tiled, row-major order iterator and a serial invocation of it

Locality Control


Role: permit tuning for locality and affinity

```
var i: int;  
on Locales[1] {  
  var j: int;  
  forall loc in Locales {  
    on loc {  
      var k: int;  
      ...  
    }
```

on-clauses migrate the current task to the specified locale (compute node)

variables are declared on the locale where the current task is running

within this scope, i and j may be referenced; the compiler and runtime manage communication



Sample mapping of tasks and their variables to locales

Next Steps

- Performance Optimizations
- Accelerator Support (GPUs, MIC)
- Improved Interoperability
- Feature Improvements
- Research Efforts
- Outreach

A Team Effort — Join Us!

