

# Chapel: Data Parallelism

# Data vs. Task Parallelism (Our Terminology)

## Data Parallelism:

- parallelism is driven by collections of data
  - data aggregates (arrays)
  - sets of indices (ranges, domains)
  - other user-defined collections
- e.g., “for all elements in array A ...”

## Task Parallelism:

- parallelism is expressed using distinct tasks
- e.g., “create a task to do foo() while another does bar()”

*(Of course, data parallelism is executed using tasks and task parallelism typically operates on data, so the line can get fuzzy at times...)*

# "Hello World" in Chapel: a Data Parallel Version

- Data Parallel Hello World

```
config const numIters = 100000;

forall i in 1..numIters do
    writeln("Hello, world! ",
            "from iteration ", i, " of ", numIters);
```

# Outline

- Domains and Arrays
  - Rectangular Domains and Arrays
  - Iterations and Operations
- Other Domain Types
- Reductions


# Domains

***Domain:*** A first-class index set

- Fundamental Chapel concept for data parallelism
- A generalization of ZPL's *region* concept
- Domains may optionally be distributed

# Sample Domains

```
config const m = 4, n = 8;  
  
var D: domain(2) = [1..m, 1..n];
```



*D*

# Sample Domains


```

config const m = 4, n = 8;

var D: domain(2) = [1..m, 1..n];

var InnerD: subdomain(D) = [2..m-1, 2..n-1];

```



# Domains Define Arrays

- Syntax


```
array-type:
  [ domain-expr ] elt-type
```

- Semantics


- Stores element for each index in *domain-expr*

- Example

```
var A, B: [D] real;
```



**A**



**B**

- Revisited example

```
var A: [1..3] int; // creates anonymous domain [1..3]
```



# Domain Iteration

- For loops (discussed already)
  - Execute loop body once per domain index, serially

```
for i in InnerD do ...
```

	1	2	3	4	5	6	
	7	8	9	10	11	12	

- Forall loops

- Executes loop body once per domain index, in parallel
- Loop must be *serializable* (executable by one task)

```
forall i in InnerD do ...
```

	.	.	.	.	.	.	
	.	.	.	.	.	.	

- Loop variables take on **const** domain index values

# Other Forall Loops

Forall loops also support...

- A shorthand notation:

```
[ (i,j) in D ] A(i,j) = i + j/10.0;
```

- Expression-based forms:

```
A = forall (i,j) in D do i + j/10.0;
```

```
A = [ (i,j) in D ] i + j/10.0;
```

1.1	1.2	1.3	1.4	1.5	1.6	1.7	1.8
2.1	2.2	2.3	2.4	2.5	2.6	2.7	2.8
3.1	3.2	3.3	3.4	3.5	3.6	3.7	3.8
4.1	4.2	4.3	4.4	4.5	4.6	4.7	4.8

**A**

# Domain Algebra

Domain values support...





- Methods for creating new domains

```
var D2 = InnerD.expand(1,0);
```

```
var D3 = InnerD.translate(0,1);
```

- Intersection via Slicing


```
var D4 = D2[D3];
```



# Array Slicing/Sub-Arrays

Indexing into arrays with domain values results in a sub-array expression

```
A[InnerD] = B[InnerD.translate(0,1)];
```




# Array Reallocation


Reassigning a domain logically reallocates its arrays

- values are preserved for common indices

```
D = [1..2*m, 1..2*n];
```



**A**



**B**

# Array Iteration

- Array expressions also support for and forall loops

```
for a in A[InnerD] do ...
```

	1	2	3	4	5	6	
	7	8	9	10	11	12	

```
forall a in A[InnerD] do ...
```

	.	.	.	.	.	.	
	.	.	.	.	.	.	

- Array loop variables refer to array values (modifiable)

```
forall (a, (i,j)) in (A, D) do a = i + j/10.0;
```

1.1	1.2	1.3	1.4	1.5	1.6	1.7	1.8
2.1	2.2	2.3	2.4	2.5	2.6	2.7	2.8
3.1	3.2	3.3	3.4	3.5	3.6	3.7	3.8
4.1	4.2	4.3	4.4	4.5	4.6	4.7	4.8

# Array Indexing

- Arrays can be indexed using variables of their domain's index type (e.g., tuples) or lists of integers

```
var i = 1, j = 2;  
var ij = (i, j);  
  
A[ij] = 1.0;  
A[i, j] = 2.0;
```

- Array indexing can use either parentheses or brackets

```
A(ij) = 3.0;  
A(i, j) = 4.0;
```

# Promoted Functions and Operators

Functions/operators expecting scalars can also take...  
 ...arrays, causing each element to be passed in

```
sin(A)
2*A
```

≈

```
forall a in A do sin(a)
forall a in A do 2*a
```

...domains, causing each index to be passed in

```
foo(Sparse)
```

≈

```
forall i in Sparse do foo(i)
```



# Where is the Parallelism?

- forall loops are implemented using multiple tasks
  - details depend on what is being iterated over
- so are operations that are equivalent to forall loops
  - promoted operators/functions, whole array assignment, ...
- many times, this parallelism can seem invisible
  - for this reason, Chapel's data parallelism can be considered *implicitly parallel*
  - it also tends to make the data parallel features easier to use and less likely to result in bugs as compared to explicit tasks


# Outline

- Domains and Arrays
- Other Domain Types
- Reductions
- NAS MG Stencil Revisited


# Chapel Domain Types

Chapel supports several domain types...


```
var OceanSpace = [0..#lat, 0..#long],
    AirSpace = OceanSpace by (2,4),
    IceSpace: sparse subdomain(OceanSpace) = genCaps();
```



*dense*




*strided*




*sparse*

*unstructured*



*associative*








```
var Vertices: domain(opaque) = ...,    People: domain(string) = ...;
```

# Chapel Array Types

All domain types can be used to declare arrays...

```
var Ocean: [OceanSpace] real,  
    Air: [AirSpace] real,  
    IceCaps[IceSpace] real;
```








```
var Weight: [Vertices] real,
```

```
Age: [People] int;
```

# Iteration

...to iterate over index sets...

```
forall ij in AirSpace do
    Ocean(ij) += IceCaps(ij);
```








```
forall v in Vertices do
    Weight(v) = numEdges(v);
```

```
forall p in People do
    Age(p) += 1;
```

# Slicing

...to slice arrays...

```
Ocean[AirSpace] += IceCaps[AirSpace];
```



...Vertices[Interior]...





...People[Interns]...

# Reallocation

...and to reallocate arrays

```
AirSpace = OceanSpace by (2,2);
```

```
IceSpace += genEquator();
```



	"steve"
	"lee"
	"sung"
	"david"
	"jacob"
	"albert"
	"brad"
	"vass"

```
newnode = Vertices.create();    People += "vass";
```

# Associative Domains and Arrays by Example

```

var Presidents: domain(string) =
    ("George", "John", "Thomas",
     "James", "Andrew", "Martin");

Presidents += "William";

var Age: [Presidents] int,
    Birthday: [Presidents] string;

Birthday("George") = "Feb 22";

forall president in President do
    if Birthday(president) == today then
        Age(president) += 1;

```

George
John
Thomas
James
Andrew
Martin
William

***Presidents***

Feb 22
Oct 30
Apr 13
Mar 16
Mar 15
Dec 5
Feb 9

***Birthday***

277
274
266
251
242
227
236

***Age***



# Outline

- Domains and Arrays
- Other Domain Types
- Reductions
- NAS MG Stencil Revisited

# Reductions

- Syntax

```
reduce-expr:
  reduce-op reduce iterator-expr
```

- Semantics

- Combines argument values using *reduce-op*
- *Reduce-op* may be built-in or user-defined

- Examples

```
total = + reduce A;
bigDiff = max reduce [i in InnerD] abs (A(i)-B(i));
(minVal, minLoc) = minloc reduce (A, D);
```

# Reduction and Scan Operators


- Built-in
  - +, \*, &&, ||, &, |, ^, min, max
  - minloc, maxloc
    - Takes a tuple of values and indices
    - Generates a tuple of the min/max value and its index
- User-defined
  - Defined via a class that supplies a set of methods
  - Compiler generates code that calls these methods
  - Based on:
 

S. J. Deitz, D. Callahan, B. L. Chamberlain, and L. Snyder. *Global-view abstractions for user-defined reductions and scans*. In Proceedings of the Eleventh ACM SIGPLAN Symposium on Principles and Practices of Parallel Programming, 2006.

# Outline

- Domains and Arrays
- Other Domain Types
- Reductions
- NAS MG Stencil Example

# Revisiting the *rprj3* Stencil from NAS MG




# NAS MG Stencil in Chapel Revisited

```


proc rprj3(S: [?SD], R: [?RD]) {
  const Stencil = [-1..1, -1..1, -1..1],
    W: [0..3] real = (0.5, 0.25, 0.125, 0.0625),
    W3D = [(i,j,k) in Stencil] W[(i!=0) + (j!=0) + (k!=0)];

  forall ijk in SD do
    S[ijk] = + reduce [offset in Stencil]
      (W3D[offset] * R[ijk + RD.stride*offset]);
}

```



S



R

# Data Parallelism: Status

- Most features implemented and working correctly
- Regular and irregular domains/arrays generating parallelism
- Scalar performance lacking in some cases (particularly higher-dimensional operations)
- Implementation of unstructured domains/arrays is correct but inefficient

# Future Directions

- Gain more experience with unstructured (graph-based) domains and arrays



# Questions?