

# Chapel: Status, HPCC, and SSICA #2

---

Sung-Eun Choi and Steve Deitz  
Cray Inc.

# Outline

- Status and Collaborations
  - About Chapel v1.1
  - Implementation status
  - External collaborations
- HPCC Benchmarks in Chapel as presented at SC '09
- SSCA #2 in Chapel

# Chapel Version 1.1

- Features
  - Open source at <http://sourceforge.net/projects/chapel/>
  - Distributed under the BSD Open Source license
  - Ported to Linux/Unix, Mac, Cygwin
- Contents
  - Compiler, runtime, standard modules, third-party libraries
  - Language spec, quick reference, numerous examples
- Highlights
  - Most data-parallel operations execute in parallel
  - Improved control of data parallelism
  - Completed Block and Cyclic distributions

# Implementation Status

- Language Basics
  - No support for inheritance from multiple or generic classes
  - Incomplete support for user-defined constructors
  - Incomplete support for sparse arrays and domains
  - Unchecked support for index types and sub-domains
  - No support for skyline arrays
  - No constant checking for domains, arrays, fields
  - Several internal memory leaks
- Task Parallelism
  - No support for atomic statements
  - Memory consistency model is not guaranteed

# Implementation Status

- Locality and Affinity
  - String assignment across locales is by reference
- Data Parallelism
  - Promoted functions/operators do not preserve shape
  - User-defined reductions are undocumented and in flux
  - No partial scans or reductions
  - Some data parallel statements are serialized
- Distributions and Layouts
  - Distributions are limited to Block and Cyclic
  - User-defined domain maps are undocumented and in flux

# Collaborations I

- **Notre Dame/ORNL** (Peter Kogge, Srinivas Sridharan, Jeff Vetter)  
Asynchronous software transactional memory over distributed memory
- **UIUC** (David Padua, Albert Sidelnik, Maria Garzaran)  
Chapel for hybrid CPU-GPU computing
- **BSC/UPC** (Alex Duran)  
Chapel over Nanos++ user-level tasking
- **U. Malaga** (Rafa Asenio, Maria Gonzales, Rafael Larossa)  
Parallel file I/O for arrays
- **OSU** (Gagan Agrawal, Bin Ren)  
User-defined reductions over FREERIDE for data intensive computing

# Collaborations II

- **U. Colorado** (Jeremy Siek, Jonathan Turner)  
Interfaces and modular generics for Chapel
- **PNNL/CASS-MT** (John Feo, Daniel Chavarria)  
Hybrid computing in Chapel; Cray XMT performance tuning; ARMCI port
- **ORNL** (David Bernholdt *et al.*, Steve Poole *et al.*)  
Code studies – Fock matrices, MADNESS, Sweep3D, coupled models, ...
- **Berkeley** (Dan Bonachea, Paul Hargrove *et al.*)  
Efficient GASNet support for Chapel; collective communication
- **U. Oregon/Paratools Inc.** (Sameer Shende)  
Performance analysis with Tau

# Outline

- Status and Collaborations
- HPCC Benchmarks in Chapel as presented at SC '09
  - Code updated for Chapel v1.1
- SSCA #2 in Chapel




# Highlights

- Global STREAM Triad 10.8 TB/s (6.4x over 2008)
  - Executed on 2048 nodes (up from 512 nodes in 2008)
  - Better scaling by eliminating extra communication
- EP STREAM Triad 12.2 TB/s
  - More similar to EP STREAM reference version
- Random Access 0.122 GUP/s (111x over 2008)
  - Executed on 2048 nodes (up from 64 nodes in 2008)
  - Optimized remote forks + better scaling as with STREAM
- A distributed-memory implementation of FFT
- A demonstration of portability
  - Cray XT4, Cray CX1, IBM pSeries 575, SGI Altix

# Chapel Implementation Characteristics

## Code Size Summary (Source Lines of Code)




# FFT and HPL in a Nutshell

- FFT
  - Uses both Block and Cyclic distributions
  - Butterfly-patterned accesses are completely local
    - Communication with nearby neighbors is local with Block
    - Communication with far off neighbors is local with Cyclic
  - Executes on distributed memory, but is slow
- HPL
  - Implementation is ready for BlockCyclic distribution
  - Executes on single locale only, but is multi-threaded


# Global STREAM Triad in Chapel (Excerpts)

```
const BlockDist = new dmap(new Block([1..m]));
```



```
const ProblemSpace:
```

```
    domain(1,int(64)) dmapped BlockDist = [1..m];
```



```
var A, B, C: [ProblemSpace] real;
```



```
forall (a,b,c) in (A,B,C) do
```

```
    a = b + alpha * c;
```


# EP STREAM Triad in Chapel (Excerpts)

```
coforall loc in Locales do on loc {
```



```
local {
```

```
var A, B, C: [1..m] real;
```



```
forall (a,b,c) in (A,B,C) do
```

```
  a = b + alpha * c;
```

```
}
```

```
}
```


# Experimental Setup

Machine Characteristics	
Model	Cray XT4
Location	ORNL
Nodes	7832
Processor	2.1 GHz Quadcore AMD Opteron
Memory	8 GB per node

Benchmark Parameters	
STREAM Triad Memory	Least value greater than 25% of memory
Random Access Memory	Least power of two greater than 25% of memory
Random Access Updates	$2^{n-10}$ for memory equal to $2^n$


# STREAM Triad Performance

**Performance of HPCC STREAM Triad (Cray XT4)**



# STREAM Triad Performance

Performance of HPCC STREAM Triad (Cray XT4)





# Global Random Access in Chapel (Excerpts)

```
const TableDist = new dmap(new Block([0..m-1])),  
    UpdateDist = new dmap(new Block([0..N_U-1]));
```

```
const TableSpace: domain ... dmapped TableDist = ...,  
    Updates: domain ... dmapped UpdateDist = ...;
```

```
var T: [TableSpace] uint(64);
```


```
forall ( ,r) in (Updates,RAStream()) do  
    on TableDist.idxToLocale(r & indexMask) {  
        const myR = r;  
        local T(myR & indexMask) ^= myR;  
    }
```

## More elegant on-block


```
on T(r&indexMask) do  
    T(r&indexMask) ^= r;
```

# Random Access Performance

**Performance of HPCC Random Access (Cray XT4)**




# Random Access Efficiency on 32+ Nodes

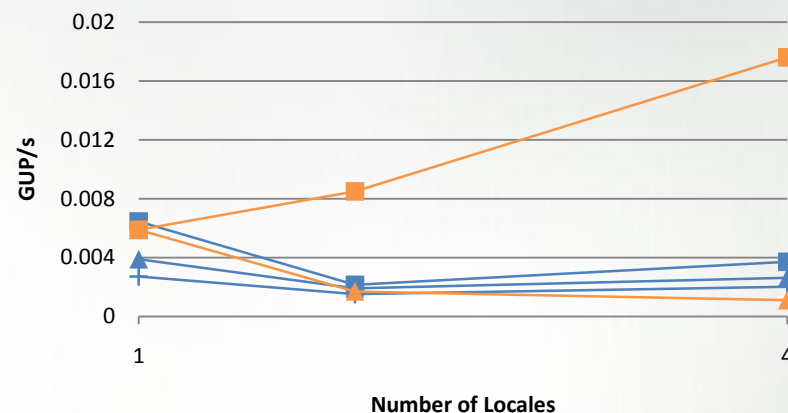


# Portability Results


Performance of HPCC STREAM Triad (Cray CX1)



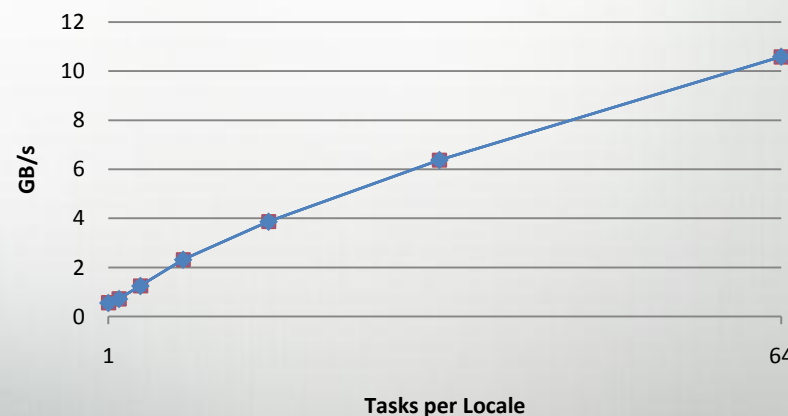
Performance of HPCC Random Access (Cray CX1)



Performance of HPCC STREAM Triad (IBM pSeries 575)



Performance of HPCC STREAM Triad (SGI Altix)




# Outline

- Status and Collaborations
- HPCC Benchmarks in Chapel as presented at SC '09
- SSCA #2 in Chapel


## SSCA #2 Kernel 2

Given a set of heavy edges *HeavyEdges* in directed graph  $G$ , find sub-graphs of outgoing paths with  $length \leq maxPathLength$



# SSCA #2 Kernel 2


Given a set of heavy edges *HeavyEdges* in directed graph *G*, find sub-graphs of outgoing paths with  $length \leq maxPathLength$



*maxPathLength = 0*

# SSCA #2 Kernel 2

Given a set of heavy edges *HeavyEdges* in directed graph *G*, find sub-graphs of outgoing paths with  $length \leq maxPathLength$




$maxPathLength = 0$      $maxPathLength = 1$



# SSCA #2 Kernel 2

Given a set of heavy edges *HeavyEdges* in directed graph *G*, find sub-graphs of outgoing paths with  $length \leq maxPathLength$



*maxPathLength* = 0    *maxPathLength* = 1    *maxPathLength* = 2

# SSCA #2 Kernel 2 (Code Courtesy of John Lewis)

```

def rootedHeavySubgraphs (
    G,
    type vertexSet;
    HeavyEdges      : domain,
    HeavyEdgeSubG    : [],
    in maxPathLength: int ) {
forall (e, subgraph) in
    (HeavyEdges, HeavyEdgeSubG) {
    const (x,y) = e;
    var ActiveLevel: vertexSet;

    ActiveLevel += y;

    subgraph.edges += e;
    subgraph.nodes += x;
    subgraph.nodes += y;
  }
}

```

```

for pathLength in 1..maxPathLength {
    var NextLevel: vertexSet;
    forall v in ActiveLevel do
        forall w in G.Neighbors(v) do
            atomic {
                if !subgraph.nodes.member(w) {
                    NextLevel += w;
                    subgraph.nodes += w;
                    subgraph.edges += (v, w);
                }
            }

    if (pathLength < maxPathLength) then
        ActiveLevel = NextLevel;
  }
}

```

# SSCA #2 Kernel 2 (Code Courtesy of John Lewis)

```
def rootedHeavySubgraphs (
    G,
    type vertexSet;
    HeavyEdges      : domain,
    HeavyEdgeSubG    : [1

    for pathLength in 1..maxPathLength {
        var NextLevel: vertexSet;
        forall v in ActiveLevel do
            forall w in G.Neighbors(v) do
```

## Generic Implementation of Graph G

**G.Vertices:** A domain whose indices represent the vertices

- For toroidal graphs, a domain( $d$ ), so vertices are  $d$ -tuples
- For other graphs, a domain(1), so vertices are integers

**G.Neighbors:** An array over G.Vertices

- For toroidal graphs, a fixed-size array over the domain  $[1..2*d]$
- For other graphs...
  - ...an associative domain with indices of type `index(G.vertices)`
  - ...a sparse subdomain of G.Vertices

```
.nodes.member(w) {
    += w;
    nodes += w;
    edges += (v, w);

    maxPathLength) then
    xtLevel;
```

*This kernel and the others are generic w.r.t. these decisions!*

# SSCA #2 Kernel 2 (Code Courtesy of John Lewis)

```
def rootedHeavySubgraphs (
    G,
    type vertexSet;
```

```
for pathLength in 1..maxPathLength {
    var NextLevel: vertexSet;
    forall v in ActiveLevel do
        forall w in G.Neighbors(v) do
```

## Generic with respect to vertex sets

**vertexSet:** A type argument specifying how to represent vertex subsets

### Requirements:

- Parallel iteration
- Ability to add members, test for membership

### Options:

- An associative domain over vertices  
**domain** (**index** (G.vertices) )
- A sparse subdomain of the vertices  
**sparse subdomain** (G.vertices)

```
atomic {
    if !subgraph.nodes.member(w) {
        NextLevel += w;
        subgraph.nodes += w;
        subgraph.edges += (v, w);
    }
}
```

```
if (pathLength < maxPathLength) then
    ActiveLevel = NextLevel;
```

# Questions?

- Status and Collaborations
  - About Chapel v1.1
  - Implementation status
  - External collaborations
- HPCC Benchmarks in Chapel as presented at SC '09
- SSCA #2 in Chapel