

Chapel: Data Parallelism

Outline

- Domains and Arrays
 - Overview
 - Arithmetic
- Other Domain Types
- Data Parallel Operations
- NAS MG Stencil Revisited

Domains


- A first-class index set
 - Specifies size and shape of arrays
 - Supports iteration, array operations
 - Potentially distributed across locales
- Three main classes
 - Arithmetic—indices are Cartesian tuples
 - Associative—indices are hash keys
 - Opaque—indices are anonymous
- Fundamental Chapel concept for data parallelism
- A generalization of ZPL's region concept

Sample Arithmetic Domains

```
config const m = 4, n = 8;

var D: domain(2) = [1..m, 1..n];

var InnerD: domain(2) = [2..m-1, 2..n-1];
```



Domains Define Arrays

- Syntax


```
array-type:
  [ domain-expr ] elt-type
```

- Semantics

- Stores element for each index in *domain-expr*

- Example

```
var A, B: [D] real;
```



- Revisited example

```
var A: [1..3] int; // creates anonymous domain [1..3]
```

Domain Iteration

- For loops (discussed already)
 - Executes loop body once per loop iteration
 - Order is serial

```
for i in InnerD do ...
```

	1	2	3	4	5	6	
	7	8	9	10	11	12	

D

- Forall loops
 - Executes loop body once per loop iteration
 - Order is parallel (must be *serializable*)

```
forall i in InnerD do ...
```

D

Other Forall Loops

Forall loops also support...

- A shorthand:

```
[ (i,j) in D] A(i,j) = i + j/10.0;
```

- An expression-based form:

```
A = forall (i,j) in D do i + j/10.0;
```

- A shorthand expression-based form:

```
A = [ (i,j) in D] i + j/10.0;
```

1.1	1.2	1.3	1.4	1.5	1.6	1.7	1.8
2.1	2.2	2.3	2.4	2.5	2.6	2.7	2.8
3.1	3.2	3.3	3.4	3.5	3.6	3.7	3.8
4.1	4.2	4.3	4.4	4.5	4.6	4.7	4.8

A


Data Parallelism Configuration Constants

- **--dataParTasksPerLocale=#**
 - Specify # of tasks to execute forall loops
 - Default: number of cores (*in current implementation*)
- **--dataParIgnoreRunningTasks=[true | false]**
 - If false, reduce # of forall tasks by # of running tasks
 - Default: true (*in current implementation*)
- **--dataParMinGranularity=#**
 - If > 0, reduce # of forall tasks if any task has fewer iterations
 - Default: 1 (*in current implementation*)

Other Domain Functionality


- Domain methods (exterior, interior, translate, ...)
- Domain slicing (intersection)
- Array slicing (sub-array references)

```
A (InnerD) = B (InnerD);
```



- Array reallocation
 - Reassign domain → change array
 - Values are preserved (new elements initialized)

```
D = [1..m+1, 1..m];
```



Array Arguments and Aliases

- Arrays are passed by reference

```
def f(A: []) { A = 0; }

f(A[InnerD]);
```

- Non-argument array alias of a slice

```
var AA => A(InnerD);
```

- Re-indexing arrays

```
def f(A: [1..n-2, 1..m-2]);

f(A[2..n-1, 2..m-1]);
```


```
var AA: [1..n-2, 1..m-2] => A[2..n-1, 2..m-1];
```

Outline


- Domains and Arrays
- Other Domain Types
 - Strided
 - Sparse
 - Associative
 - Opaque
- Data Parallel Operations
- NAS MG Stencil Revisited

The Varied Kinds of Domains


```
var Dense: domain(2) = [1..10, 1..20],
    Strided: domain(2) = Dense by (2, 4),
    Sparse: sparse subdomain(Dense) = genIndices(),
    Associative: domain(string) = readNames(),
    Opaque: domain(opaque);
```




Dense



Strided



Sparse




Opaque

George
John
Thomas
James
Andrew
Martin
William


Associative

The Varied Kinds of Arrays


```
var DenseArr: [Dense] real,  
    StridedArr: [Strided] real,  
    SparseArr: [Sparse] real,  
    AssociativeArr: [Associative] real,  
    OpaqueArr: [Opaque] real;
```




DenseArr



StridedArr



SparseArr




OpaqueArr

George
John
Thomas
James
Andrew
Martin
William


AssociativeArr

All Domains Support Iteration


```
forall (i,j) in Strided {
  DenseArr(i,j) += SparseArr(i,j);
}
```




DenseArr



Strided



SparseArr



OpaqueArr

George
John
Thomas
James
Andrew
Martin
William

AssociativeArr

(Also, all domains support slicing, reallocation, ...)

Associative Domains and Arrays by Example

```

var Presidents: domain(string) =
    ("George", "John", "Thomas",
     "James", "Andrew", "Martin");

Presidents += "William";

var Ages: [Presidents] int,
    Birthdays: [Presidents] string;

Birthdays("George") = "Feb 22";

forall president in Presidents do
    if Birthdays(president) == today then
        Ages(president) += 1;

```

George
John
Thomas
James
Andrew
Martin
William

Presidents

Feb 22
Oct 30
Apr 13
Mar 16
Mar 15
Dec 5
Feb 9

Birthdays

277
274
266
251
242
227
236

Ages

More Domain and Array Examples

- `examples/primers/slices.chpl`
- `examples/primers/sparse.chpl`
- `examples/primers/opaque.chpl`

Outline

- Domains and Arrays
- Other Domain Types
- Data Parallel Operations
 - Promotion
 - Reductions
 - Scans
- NAS MG Stencil Revisited

Data Parallel Promotion

Functions/operators expecting scalars can also take...

- Arrays, causing each element to be passed

$\sin(A)$ $2 * A$	\approx	forall a in A do $\sin(a)$ forall a in A do $2 * a$
----------------------	-----------	--

- Domains, causing each index to be passed

foo(Sparse)	\approx	forall i in Sparse do foo(i)
-------------	-----------	---

Multiple arguments can promote using either...

- Zipper promotion

pow(A, B)	\approx	forall (a,b) in (A,B) do pow(a,b)
-----------	-----------	--

- Tensor product promotion

pow[A, B]	\approx	forall (a,b) in [A,B] do pow(a,b)
-----------	-----------	--

Reductions

- Syntax

```
reduce-expr:  
  reduce-op reduce iterator-expr
```

- Semantics

- Combines iterated elements with *reduce-op*
- *Reduce-op* may be built-in or user-defined

- Examples

```
total = + reduce A;  
bigDiff = max reduce [i in InnerD] abs(A(i)-B(i));
```

Scans

- Syntax

```
scan-expr:
  scan-op scan iterator-expr
```

- Semantics

- Computes parallel prefix of *scan-op* over elements
- *Scan-op* may be any *reduce-op*

- Examples

```
var A, B, C: [1..5] int;  
A = 1; // A: 1 1 1 1 1  
B = + scan A; // B: 1 2 3 4 5  
B(3) = -B(3); // B: 1 2 -3 4 5  
C = min scan B; // C: 1 1 -3 -3 -3
```

Reduction and Scan Operators

- Built-in
 - +, *, &&, ||, &, |, ^, min, max
 - minloc, maxloc

(Generate a tuple of the min/max and its index)
- User-defined
 - Defined via a class that supplies a set of methods
 - Compiler generates code that calls these methods
 - More information:

S. J. Deitz, D. Callahan, B. L. Chamberlain, and L. Snyder. *Global-view abstractions for user-defined reductions and scans*. In Proceedings of the Eleventh ACM SIGPLAN Symposium on Principles and Practices of Parallel Programming, 2006.


Reduction Examples

- `examples/primers/reductions.chpl`

Outline

- Domains and Arrays
- Other Domain Types
- Data Parallel Operations
- NAS MG Stencil Revisited

NAS MG Stencil Revisited



NAS MG Stencil in Chapel Revisited

```

def rprj3(S, R) {
  const Stencil = [-1..1, -1..1, -1..1],
    W: [0..3] real = (0.5, 0.25, 0.125, 0.0625),
    W3D = [(i,j,k) in Stencil] W((i!=0)+(j!=0)+(k!=0));

  forall inds in S.domain do
    S(inds) =
      + reduce [offset in Stencil] (W3D(offset) *
                                     R(inds + offset*R.stride));
}
  
```

Questions?

- Domains and Arrays
 - Overview
 - Arithmetic
- Other Domain Types
 - Strided
 - Sparse
 - Associative
 - Opaque
- Data Parallel Operations
 - Promotion
 - Reductions
 - Scans
- NAS MG stencil revisited