



Hewlett Packard
Enterprise

ONE-DAY CHAPEL TUTORIAL

SESSION 3: PARALLELISM IN CHAPEL

Chapel Team

October 16, 2023

ONE DAY CHAPEL TUTORIAL

- 9-10:30: Getting started using Chapel for parallel programming
- 10:30-10:45: break
- 10:45-12:15: Chapel basics in the context of the n-body example code
- 12:15-1:15: lunch
- 1:15-2:45: Distributed and shared-memory parallelism especially w/arrays (data parallelism)
- 2:45-3:00: break
- 3:00-4:30: More parallelism including for asynchronous parallelism (task parallelism)
- 4:30-5:00: Wrap-up including gathering further questions from attendees



OUTLINE: PARALLELISM IN CHAPEL

- Recall processing files in parallel
- Data parallelism concepts and examples including multi-locale parallelism with distributions
- Domains
- Forall Loops
- Domain Distributions
- Using a Different Domain Distribution
- Implicit Communication: Remote writes/Puts and Reads/Gets
- Parallelizing a 1D heat diffusion solver (Hands On)
- Heat 2D example with CommDiagnostics (Hands On)



RECALL PROCESSING FILES IN PARALLEL

RECALL: ANALYZING MULTIPLE FILES USING PARALLELISM

 parfilekmer.chpl

parfilekmer.chpl

```
use FileSystem;
config const dir = "DataDir";
var fList = findFiles(dir);
var filenames =
    blockDist.createArray(0..<fList.size, string);
filenames = fList;

// per file word count
forall f in filenames {
    ...
    // code from kmer.chpl
    ...
}
```

```
prompt> chpl --fast parfilekmer.chpl
prompt> ./parfilekmer -nl 1
prompt> ./parfilekmer -nl 4
```

- shared and distributed-memory parallelism using 'forall'
 - in other words, parallelism within the locale/node and across locales/nodes
- a distributed array
- command line options to indicate number of locales



RECALL: BLOCK DISTRIBUTION OF ARRAY OF STRINGS

Locale 0

Locale 1

"filename1"	"filename2"	"filename3"	"filename4"	"filename5"	"filename6"	"filename7"	"filename8"
-------------	-------------	-------------	-------------	-------------	-------------	-------------	-------------

```
prompt> chpl --fast parfilekmer.chpl
prompt> ./parfilekmer -nl 2
```

- Array of strings for filenames is distributed across locales
- 'forall' will do parallelism across locales and then within each locale to take advantage of multicore



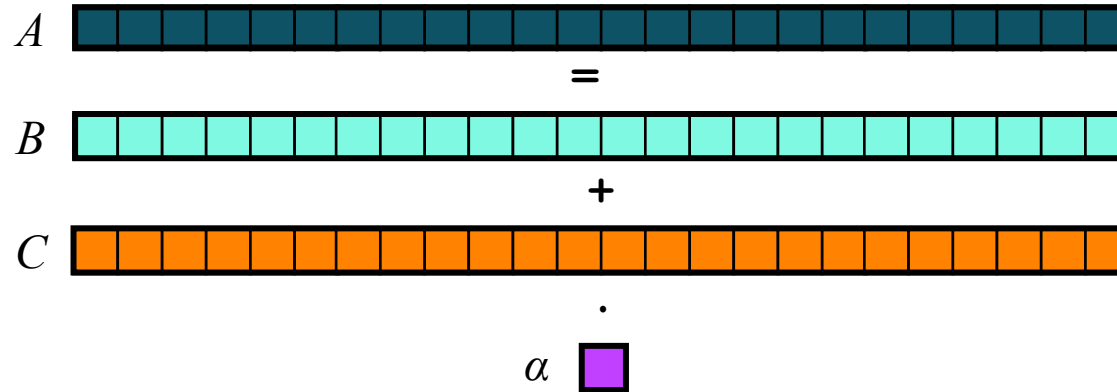
DATA PARALLELISM CONCEPTS AND EXAMPLES INCLUDING MULTI-LOCALE PARALLELISM WITH DISTRIBUTIONS

STREAM TRIAD: A PARALLEL COMPUTATION

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures:

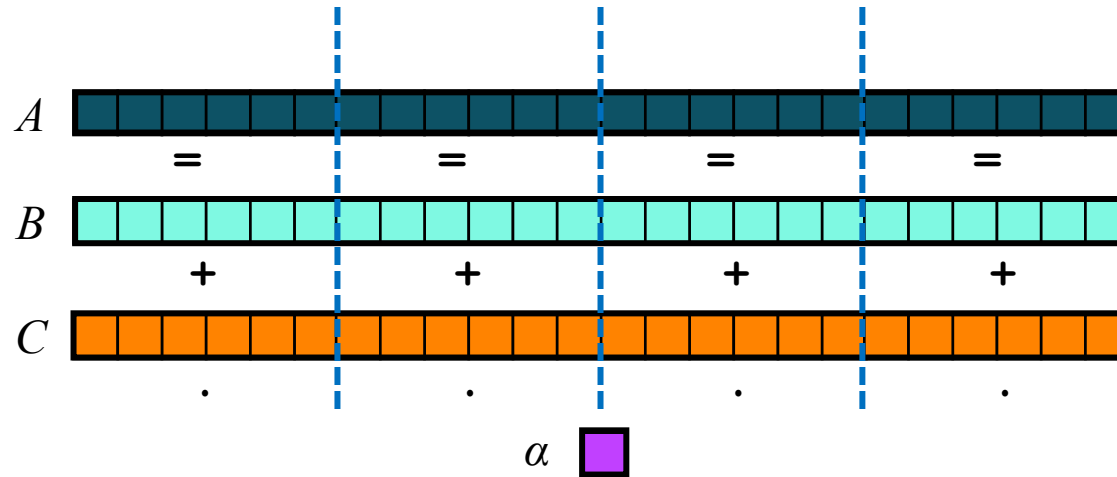


STREAM TRIAD: A PARALLEL COMPUTATION

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (shared memory / multicore):

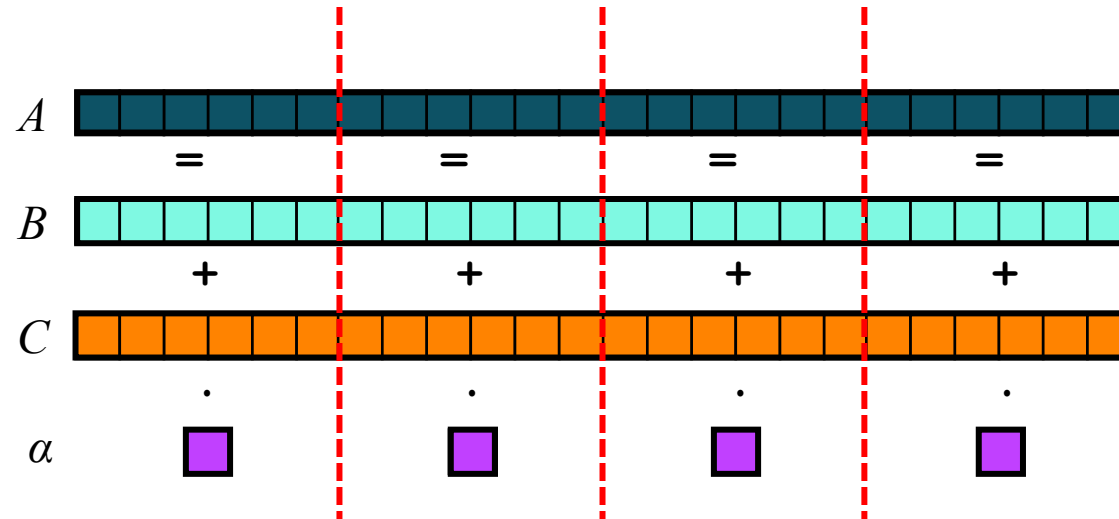


STREAM TRIAD: A PARALLEL COMPUTATION

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (distributed memory):

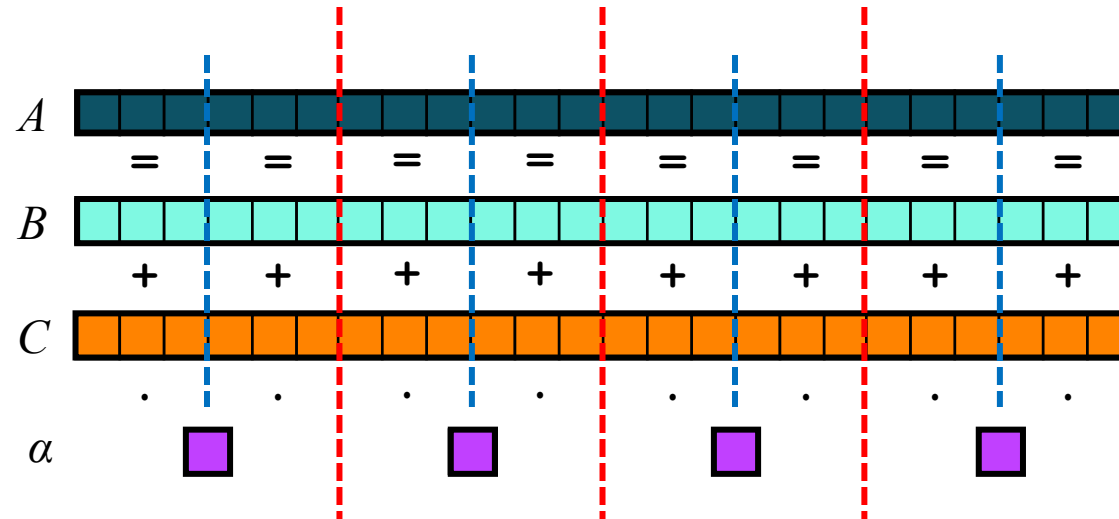


STREAM TRIAD: A PARALLEL COMPUTATION

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (distributed memory multicore):

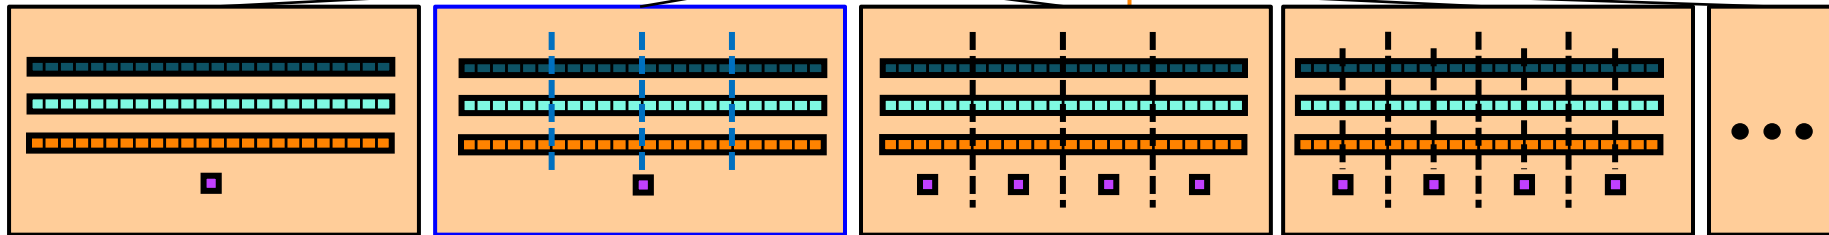


STREAM TRIAD: CHAPEL

```
use BlockDist;  
  
config const m = 1000,  
           alpha = 3.0;  
  
const ProblemSpace = blockDist.createDomain({1..m});  
  
var A, B, C: [ProblemSpace] real;  
  
B = 2.0;  
C = 1.0;  
  
A = B + alpha * C;
```

The special sauce:

How should this index set—and any arrays and computations over it—be mapped to the system?



Philosophy: Good, *top-down* language design can tease system-specific implementation details away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.

DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

```
config const n = 1000;  
var D = {1..n, 1..n};  
  
var A: [D] real;  
forall (i,j) in D with (ref A) do  
    A[i,j] = i + (j - 0.5)/n;  
writeln(A);
```


```
prompt> chpl dataParallel.chpl  
prompt> ./dataParallel -nl 1 --n=5  
1.1 1.3 1.5 1.7 1.9  
2.1 2.3 2.5 2.7 2.9  
3.1 3.3 3.5 3.7 3.9  
4.1 4.3 4.5 4.7 4.9  
5.1 5.3 5.5 5.7 5.9
```

DOMAINS

DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Domains (Index Sets)



```
config const n = 1000;  
var D = {1..n, 1..n};  
  
var A: [D] real;  
forall (i,j) in D with (ref A) do  
    A[i,j] = i + (j - 0.5)/n;  
writeln(A);
```

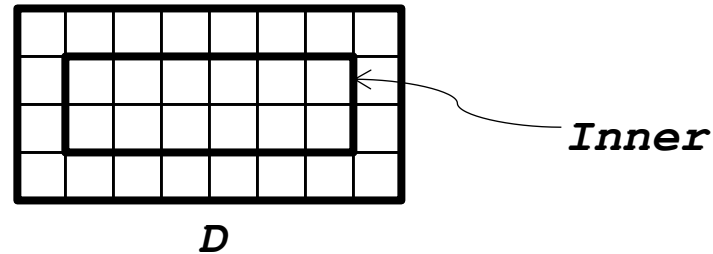
```
prompt> chpl dataParallel.chpl  
prompt> ./dataParallel -nl 1 --n=5  
1.1 1.3 1.5 1.7 1.9  
2.1 2.3 2.5 2.7 2.9  
3.1 3.3 3.5 3.7 3.9  
4.1 4.3 4.5 4.7 4.9  
5.1 5.3 5.5 5.7 5.9
```

DOMAINS

Domain:

- A first-class index set
- The fundamental Chapel concept for data parallelism

```
config const m = 4, n = 8;  
  
const D = {1..m, 1..n};  
const Inner = {2..m-1, 2..n-1};
```

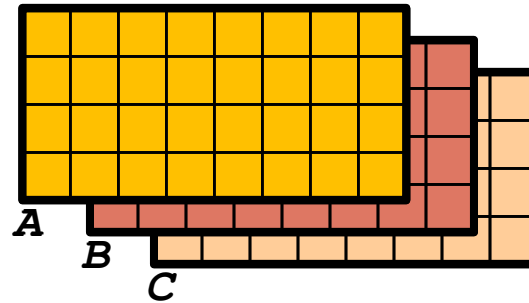


DOMAINS

Domain:

- A first-class index set
- The fundamental Chapel concept for data parallelism
- Useful for declaring arrays and computing with them

```
config const m = 4, n = 8;  
  
const D = {1..m, 1..n};  
const Inner = {2..m-1, 2..n-1};  
  
var A, B, C: [D] real;
```



DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Arrays



```
config const n = 1000;  
var D = {1..n, 1..n};  
  
var A: [D] real;  
forall (i,j) in D with (ref A) do  
    A[i,j] = i + (j - 0.5)/n;  
writeln(A);
```

```
prompt> chpl dataParallel.chpl  
prompt> ./dataParallel -nl 1 --n=5  
1.1 1.3 1.5 1.7 1.9  
2.1 2.3 2.5 2.7 2.9  
3.1 3.3 3.5 3.7 3.9  
4.1 4.3 4.5 4.7 4.9  
5.1 5.3 5.5 5.7 5.9
```

FORALL LOOPS

DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Data-Parallel Forall Loops

```
config const n = 1000;  
var D = {1..n, 1..n};  
  
var A: [D] real;  
forall (i,j) in D with (ref A) do  
    A[i,j] = i + (j - 0.5)/n;  
writeln(A);
```

```
prompt> chpl dataParallel.chpl  
prompt> ./dataParallel -nl 1 --n=5  
1.1 1.3 1.5 1.7 1.9  
2.1 2.3 2.5 2.7 2.9  
3.1 3.3 3.5 3.7 3.9  
4.1 4.3 4.5 4.7 4.9  
5.1 5.3 5.5 5.7 5.9
```

FORALL LOOPS

Forall loops: Central concept for data parallel computation

- Like for-loops, but parallel
- Implementation details determined by iterand (e.g., D below)
 - specifies number of tasks, which tasks run which iterations, ...
 - in practice, typically uses a number of tasks appropriate for target HW

```
forall (i,j) in D with (ref A) do  
    A[i,j] = i + j/10.0;
```

Forall loops assert...

- ...parallel safety:** OK to execute iterations simultaneously
- ...order independence:** iterations could occur in any order
- ...serializability:** all iterations could be executed by one task
 - e.g., can't have synchronization dependences between iterations

1.1	1.2	1.3	1.4	1.5	1.6	1.7	1.8
2.1	2.2	2.3	2.4	2.5	2.6	2.7	2.8
3.1	3.2	3.3	3.4	3.5	3.6	3.7	3.8
4.1	4.2	4.3	4.4	4.5	4.6	4.7	4.8

COMPARISON OF LOOPS: FOR, FORALL, AND COFORALL

For loops: executed using one task

- use when a loop must be executed serially
- or when one task is sufficient for performance

Forall loops: typically executed using $1 < \#tasks \ll \#iters$

- use when a loop *should* be executed in parallel...
- ...but *can* legally be executed serially
- use when desired $\# tasks \ll \#$ of iterations

Coforall loops: executed using a task per iteration

- use when the loop iterations *must* be executed in parallel
- use when you want $\# tasks == \#$ of iterations
- use when each iteration has substantial work



DATA PARALLELISM, BY EXAMPLE



03-domain-distributions.chpl

This is a shared memory program

Nothing has referred to remote
locales, explicitly or implicitly

```
config const n = 1000;  
var D = {1..n, 1..n};  
  
var A: [D] real;  
forall (i,j) in D with (ref A) do  
    A[i,j] = i + (j - 0.5)/n;  
writeln(A);
```

```
prompt> chpl dataParallel.chpl  
prompt> ./dataParallel -nl 1 --n=5  
1.1 1.3 1.5 1.7 1.9  
2.1 2.3 2.5 2.7 2.9  
3.1 3.3 3.5 3.7 3.9  
4.1 4.3 4.5 4.7 4.9  
5.1 5.3 5.5 5.7 5.9
```

DOMAIN DISTRIBUTIONS

DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Domain Distribution
(Map Data Parallelism to the System)

```
use CyclicDist;
config const n = 1000;
var D = cyclicDist.createDomain({1..n, 1..n});


var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl
prompt> ./dataParallel --n=5 -nl 4
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```

DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

High-level distributed and shared
memory parallelism



```
use CyclicDist;
config const n = 1000;
var D = cyclicDist.createDomain({1..n, 1..n});

var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

Provides programmability and control

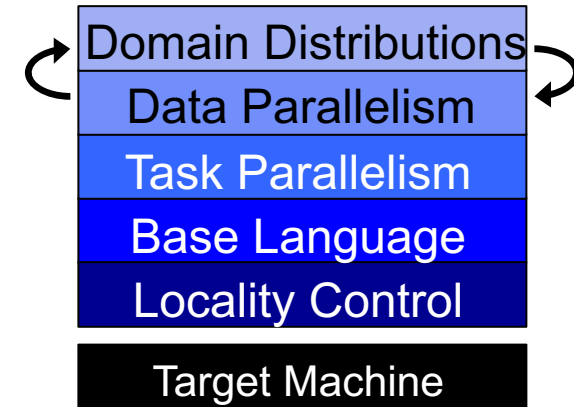
- Lowering of code is well-defined
- User can control details
- Part of Chapel's *multiresolution philosophy*...

```
prompt> chpl dataParallel.chpl
prompt> ./dataParallel --n=5 --nl 4
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```

CHAPEL'S MULTIRESOLUTION PHILOSOPHY

Multiresolution Design: Support multiple tiers of features

- higher levels for programmability, productivity
- lower levels for greater degrees of control
- build the higher-level concepts in terms of the lower
- permit users to intermix layers arbitrarily



DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Chapel's prescriptive approach:

```
forall (i,j) in D do...
```

- ⇒ invoke and inline D's
default parallel iterator
- defined by D's type /
domain distribution

default domain distribution

- create a task per local core
- block indices across tasks

```
config const n = 1000;  
var D = {1..n, 1..n};  
  
var A: [D] real;  
forall (i,j) in D with (ref A) do  
    A[i,j] = i + (j - 0.5)/n;  
writeln(A);
```

```
prompt> chpl dataParallel.chpl  
prompt> ./dataParallel --n=5 -nl 1  
1.1 1.3 1.5 1.7 1.9  
2.1 2.3 2.5 2.7 2.9  
3.1 3.3 3.5 3.7 3.9  
4.1 4.3 4.5 4.7 4.9  
5.1 5.3 5.5 5.7 5.9
```

DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Chapel's prescriptive approach:

```
forall (i,j) in D do...
```

⇒ invoke and inline D's
default parallel iterator

- defined by D's type /
domain distribution

cyclic domain distribution

on each target locale...

- create a task per core
- block local indices across tasks

```
use CyclicDist;
config const n = 1000;
var D = cyclicDist.createDomain({1..n, 1..n});

var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

```
prompt> chpl dataParallel.chpl
prompt> ./dataParallel --n=5 -nl=4
1.1 1.3 1.5 1.7 1.9
2.1 2.3 2.5 2.7 2.9
3.1 3.3 3.5 3.7 3.9
4.1 4.3 4.5 4.7 4.9
5.1 5.3 5.5 5.7 5.9
```

DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Chapel's prescriptive approach:

```
forall (i,j) in D do...
```

What if I don't like D's iteration strategy?

```
use CyclicDist;
config const n = 1000;
var D = cyclicDist.createDomain({1..n, 1..n});
var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

Write and call your own parallel iterator:

```
forall (i,j) in myParIter(D) do...
```



DISTRIBUTED DATA PARALLELISM, BY EXAMPLE

 03-domain-distributions.chpl

Chapel's prescriptive approach:

```
forall (i,j) in D do...
```

What if I don't like D's iteration strategy?

```
use CyclicDist;
config const n = 1000;
var D = cyclicDist.createDomain({1..n, 1..n});
var A: [D] real;
forall (i,j) in D with (ref A) do
    A[i,j] = i + (j - 0.5)/n;
writeln(A);
```

Write and call your own parallel iterator:

```
forall (i,j) in myParIter(D) do...
```

Or use a different domain distribution:

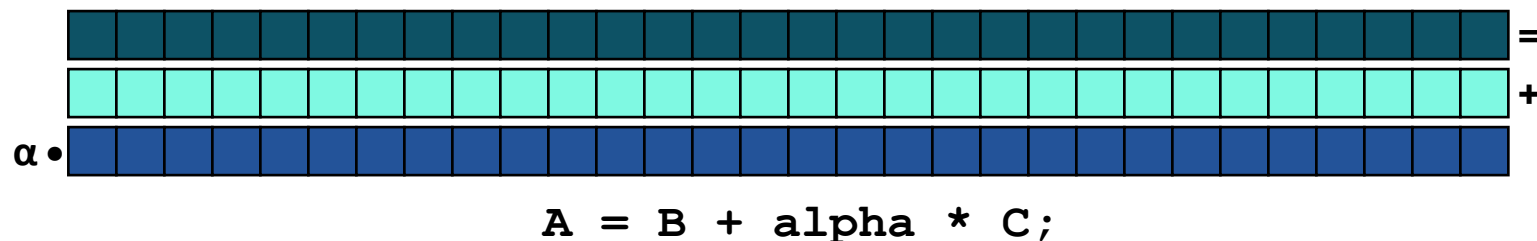
```
var D = blockDist.createDomain({1..n, 1..n});
```



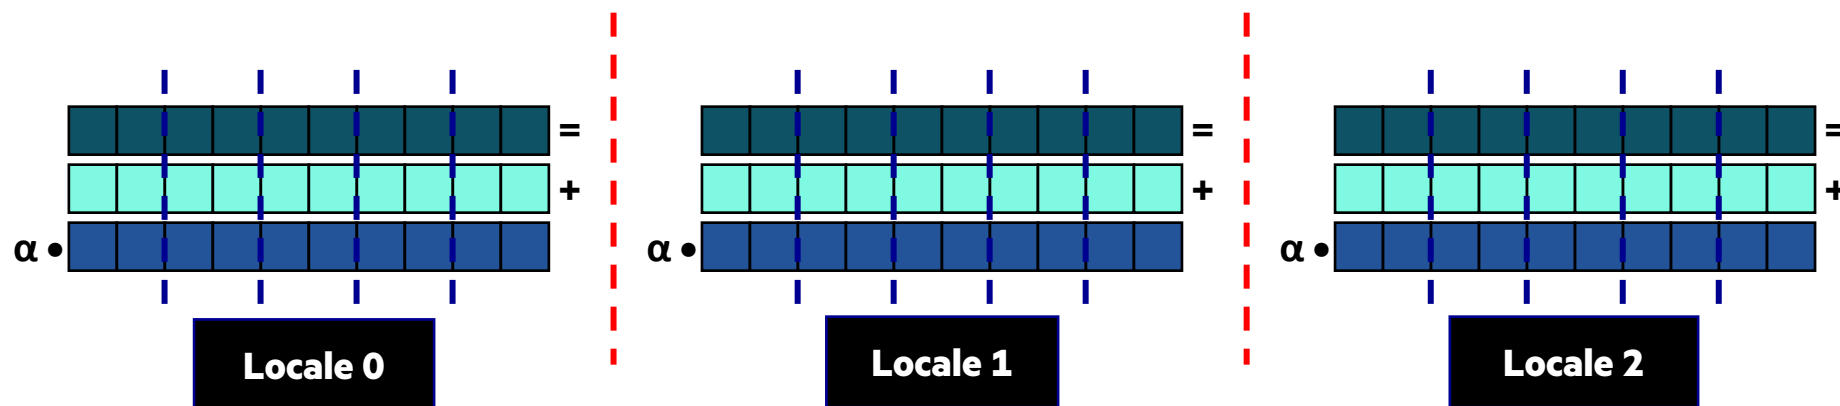
USING A DIFFERENT DOMAIN DISTRIBUTION

DOMAIN DISTRIBUTIONS: A MULTIREOLUTION FEATURE

Domain distributions are “recipes” that instruct the compiler how to map the global view of a computation...

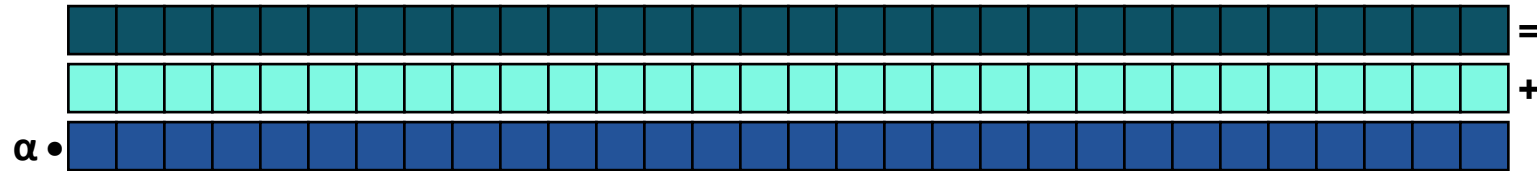


...to the target locales' memory and processors:



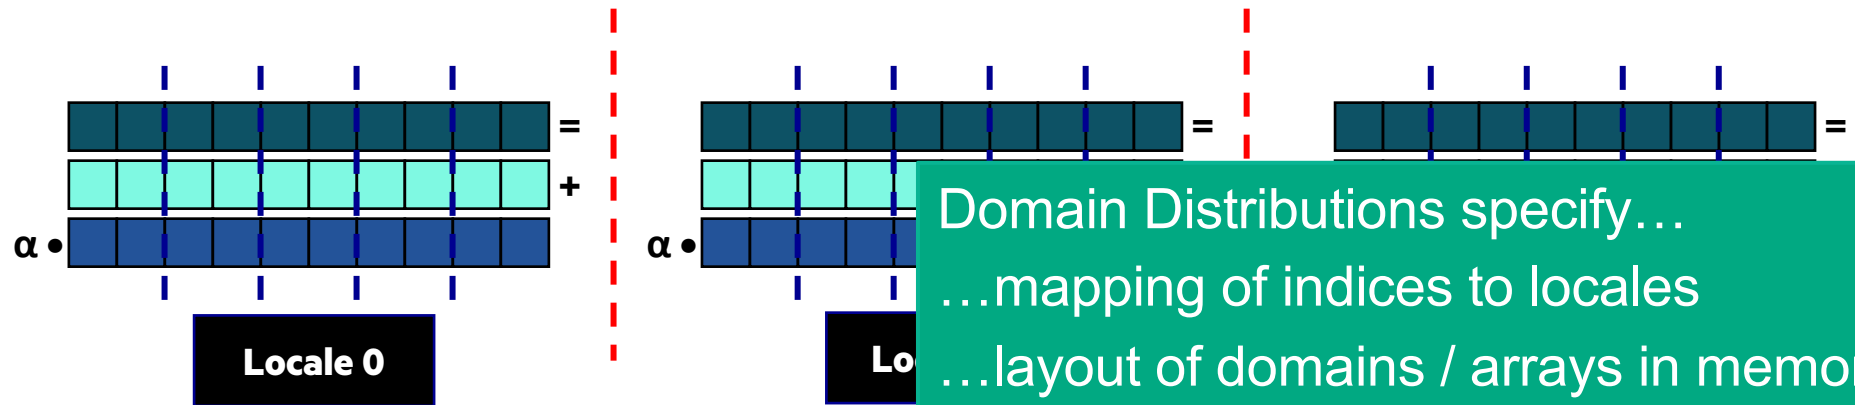
DOMAIN DISTRIBUTIONS: A MULTIREOLUTION FEATURE

Domain distributions are “recipes” that instruct the compiler how to map the global view of a computation...



$$A = B + \text{alpha} * C;$$

...to the target locales' memory and processors:



Domain Distributions specify...

...mapping of indices to locales

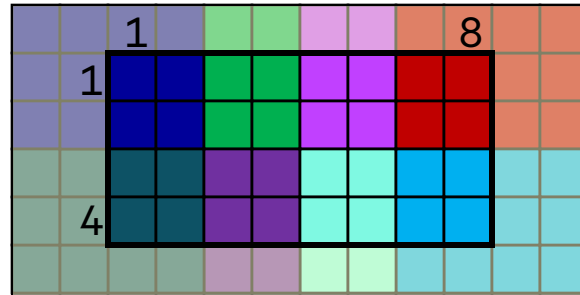
...layout of domains / arrays in memory

...parallel iteration strategies

...core operations on arrays / domains

SAMPLE DOMAIN DISTRIBUTIONS: BLOCK AND CYCLIC

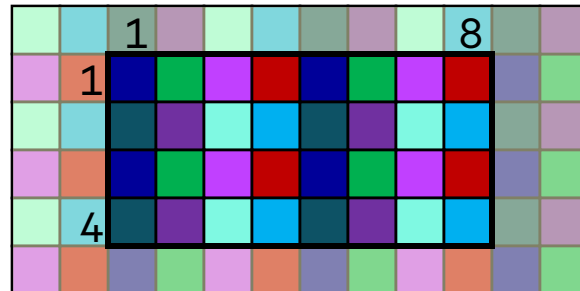
```
var Dom = blockDist.createDomain({1..4, 1..8});
```



distributed to

L0	L1	L2	L3
L4	L5	L6	L7

```
var Dom = cyclicDist.createDomain({1..4, 1..8});
```



distributed to

L0	L1	L2	L3
L4	L5	L6	L7

**IMPLICIT COMMUNICATION:
REMOTE WRITES/PUTS AND READS/GETS**

CHAPEL SUPPORTS A GLOBAL NAMESPACE WITH PUTS AND GETS

Note 1: Variables are allocated on the locale where the task is running

 03-onClause.chpl

03-onClause.chpl

```
config const verbose = false;  
var total = 0,  
    done = false;  
  
...  
  
on Locales[1] {  
    var x, y, z: int;  
    ...  
}
```

verbose false
total 0
done false

locale 0

x 0
y 0
z 0

locale 1

CHAPEL SUPPORTS A GLOBAL NAMESPACE WITH PUTS AND GETS

Note 2: Tasks can refer to lexically visible variables, whether local or remote

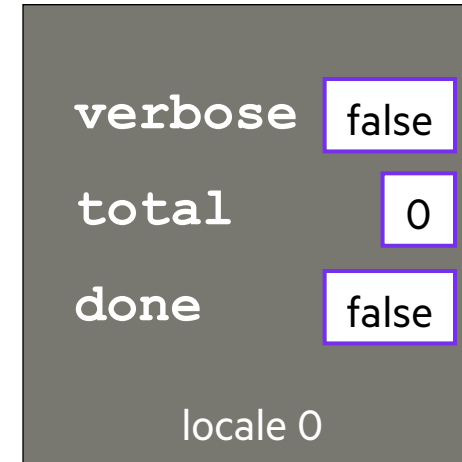
 03-onClause.chpl

03-onClause.chpl

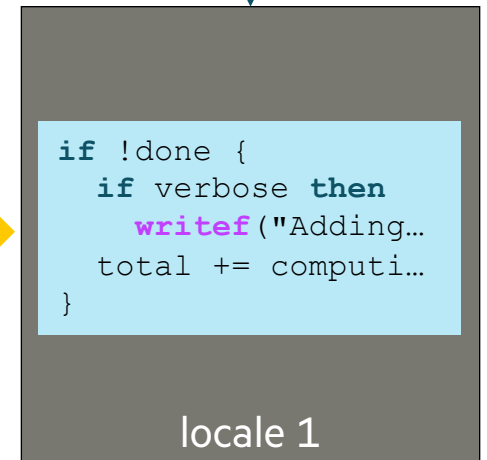
```
config const verbose = false;
var total = 0,
    done = false;

...

on Locales[1] {
  if !done {
    if verbose then
      writef("Adding locale 1's contribution");
    total += computeMyContribution();
  }
}
```



code runs on locale 1,
but refers to values
stored on locale 0



ARRAY-BASED PARALLELISM AND LOCALITY



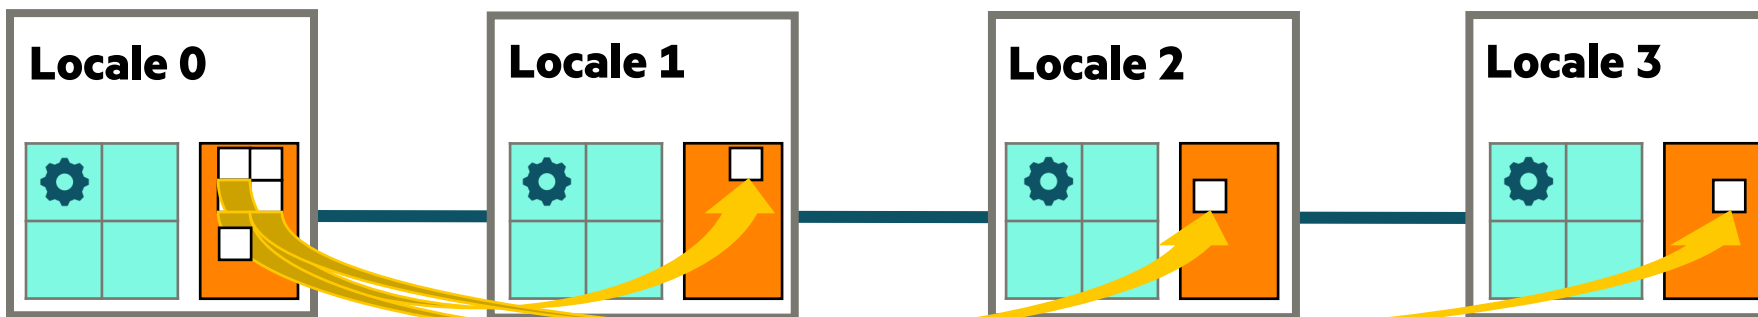
03-basics-distarr.chpl

03-basics-distarr.chpl

```
writeln("Hello from locale ", here.id);  
  
var A: [1..2, 1..2] real;  
  
use BlockDist;  
  
var D = blockDist.createDomain({1..2, 1..2});  
var B: [D] real;  
B = A;
```

Chapel also supports distributed domains (index sets) and arrays

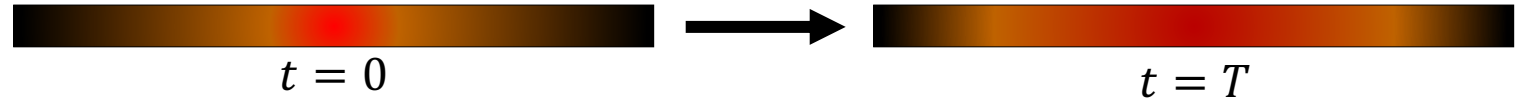
They also result in parallel distributed computation



PARALLELIZING A 1D HEAT DIFFUSION SOLVER (HANDS ON)

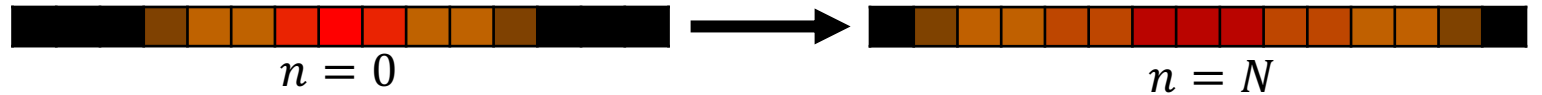
1D HEAT EQUATION EXAMPLE

Differential equation: $\frac{\partial u}{\partial t} = \alpha \frac{\partial^2 u}{\partial x^2}$



Discretized (finite difference) equation: $u_i^{n+1} = u_i^n + \alpha (u_{i-1}^n - 2u_i^n + u_{i+1}^n)$

- where $i \in \Omega \subset \mathbb{R}^1$ are discrete points in space, and $(n, n+1, \dots)$ are discrete instances in time



Finite difference algorithm:

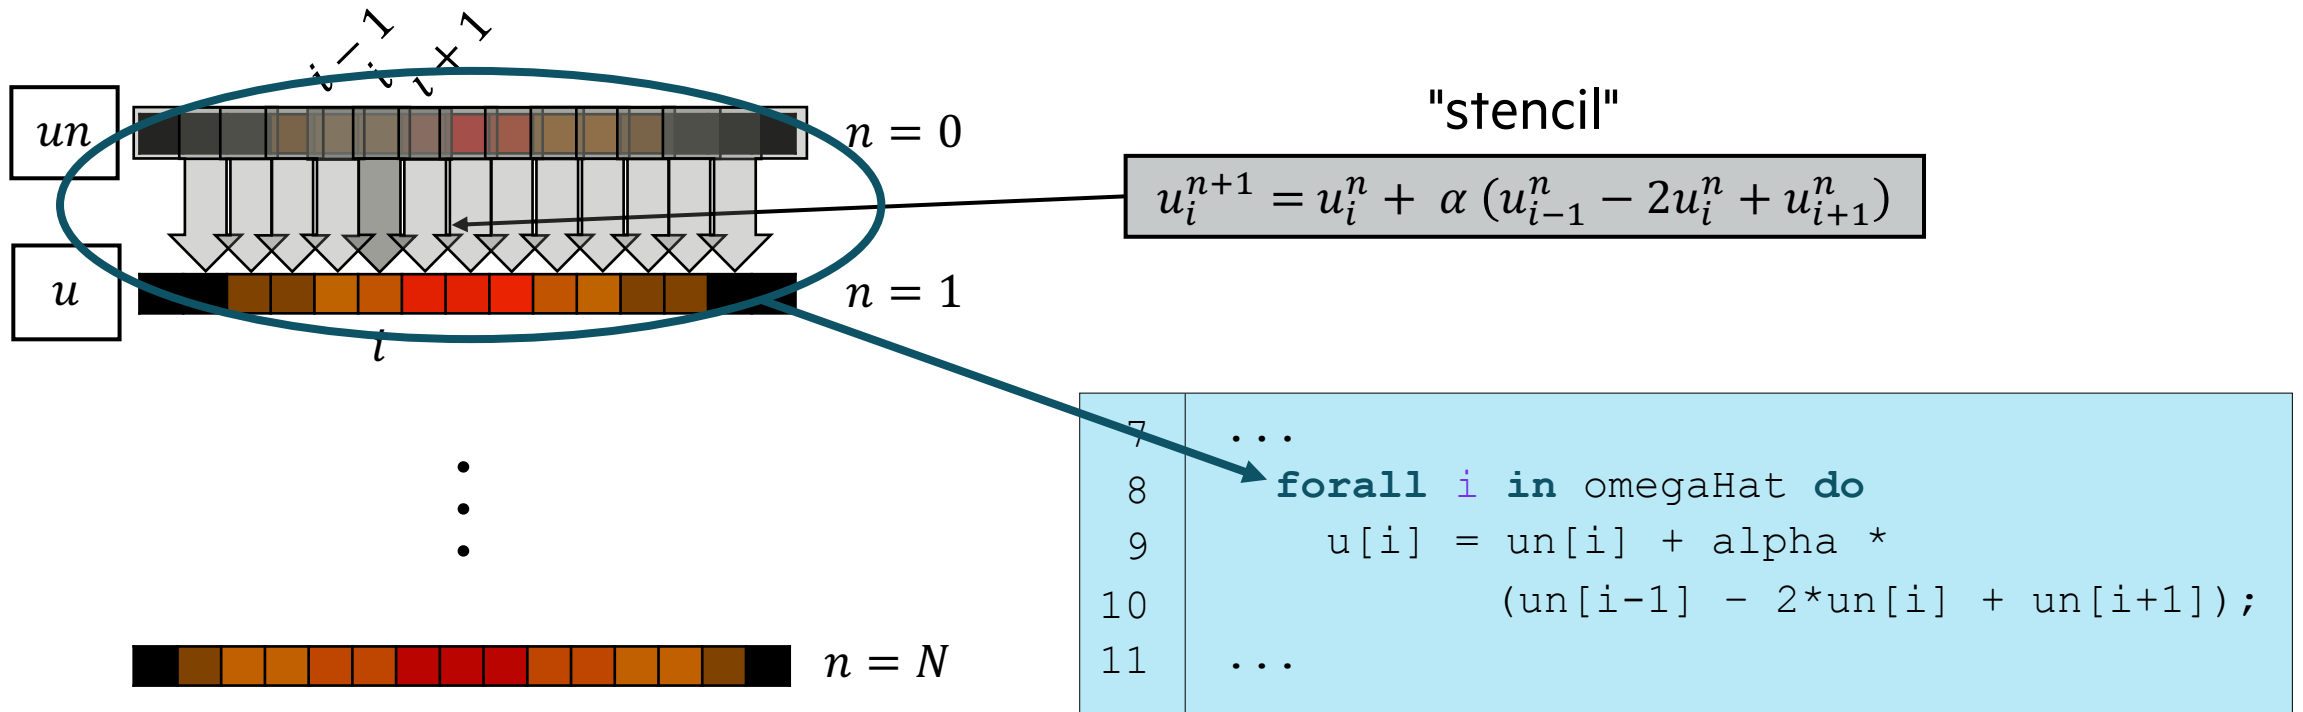
- define Ω to be a set of discrete points along the x-axis
- define $\hat{\Omega}$ over the same points, excluding the boundaries
- define an array u to over Ω
- set some initial conditions
- create a temporary copy of u , named un
- for N timesteps:
 - (1) swap u and un
 - (2) compute u in terms of un over $\hat{\Omega}$

```
1  const omega = {0..<nx},
2      omegaHat = omega.expand(-1);
3  var u: [omega] real = 1.0;
4  u[nx/4..3*nx/4] = 2.0;
5  var un = u;
6  for 1..N {
7      un <=> u;
8      forall i in omegaHat do
9          u[i] = un[i] + alpha *
10              (un[i-1] - 2*un[i] + un[i+1]);
11  }
```

1D HEAT EQUATION EXAMPLE

This pattern is often referred to as a **Stencil Computation**

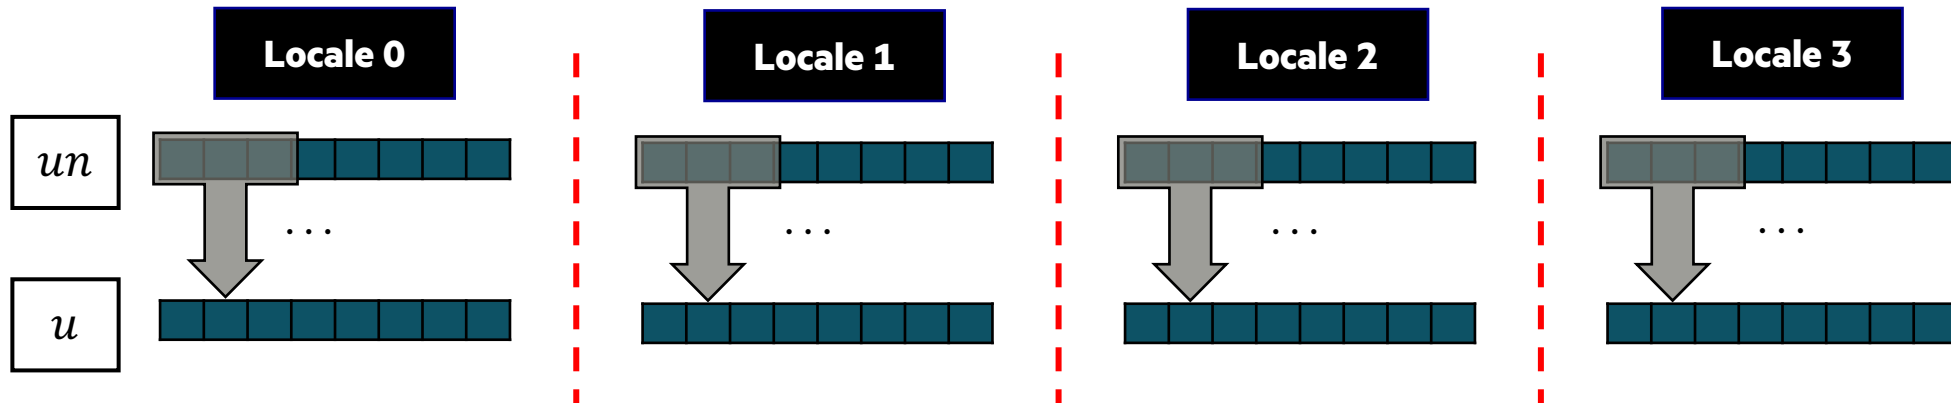
- The values in the array can be computed by applying a "stencil" to its previous state
- Note that in this case, the stencil can be applied to the entire array in parallel
 - each value in un depends strictly on values in u



HANDS ON: DISTRIBUTING THE 1D HEAT EQUATION

Imagine we want to simulate a very large domain

- We could use the Block distribution to distribute u and un across multiple locales
 - taking advantage of their memory and compute resources



Look at **heat-1D-block.chpl** and fill in the blanks to make the arrays block-distributed

Hint | Define a block-distributed domain:

```
use BlockDist;  
...  
const myBlockDom = blockDist.createDomain({1..10});
```

HANDS ON: DISTRIBUTING THE 1D HEAT EQUATION

 heat-1D-block-solution.chpl

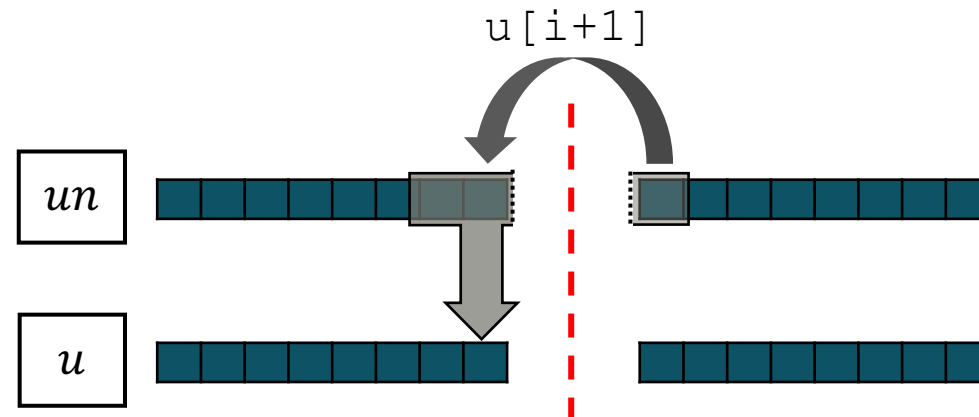
Solution: make 'omega' block-distributed:

```
omega = blockDist.createDomain({0.. $\text{nx}$ });
```

Why does this work?

- 'omegaHat' inherits 'omega's distribution
- 'u' is block-distributed
- 'un' inherits 'u's domain (and distribution)
- 'omegaHat' invokes 'blockDist's parallel/distr. iterator
 - the body of the loop is automatically split across multiple tasks on each locale
- Communication occurs automatically when a loop references a value stored on a remote locale

```
1  const omega =  
2      blockDist.createDomain({0.. $\text{nx}$ }),  
3      omegaHat = omega.expand(-1);  
4  var u: [omega] real = 1.0;  
5  u[nx/4.. $3*\text{nx}/4$ ] = 2.0;  
6  var un = u;  
7  for 1.. $N$  {  
8      un <=> u;  
9      forall i in omegaHat do  
10         u[i] = un[i] + alpha *  
11             (un[i-1] - 2*un[i] + un[i+1]);  
12 }
```



HEAT 2D EXAMPLE WITH COMMDIAGNOSTICS (HANDS ON)

2D HEAT EQUATION EXAMPLE

2D and 3D stencil codes are more common and practical

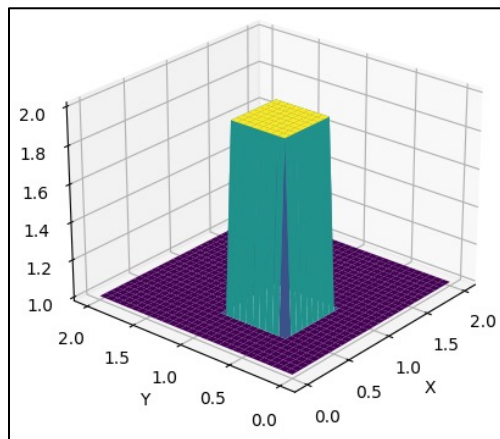
- They also present more interesting considerations for parallelization and distribution

2D heat / diffusion PDE:

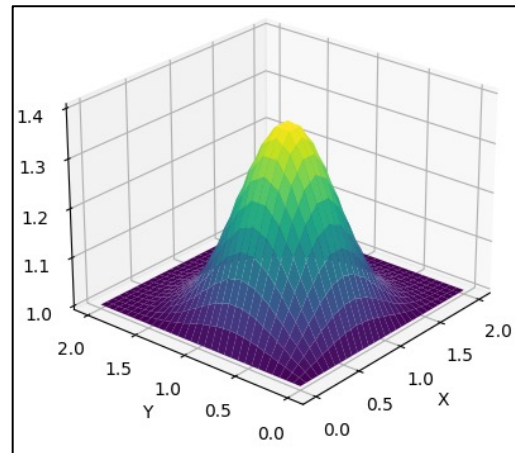
$$\frac{\partial u}{\partial t} = \alpha \Delta u = \alpha \left(\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} \right)$$

Discretized (finite-difference) form:

$$u_{i,j}^{n+1} = u_{i,j}^n + \alpha (u_{i+1,j}^n + u_{i-1,j}^n - 4u_{i,j}^n + u_{i,j+1}^n + u_{i,j-1}^n)$$



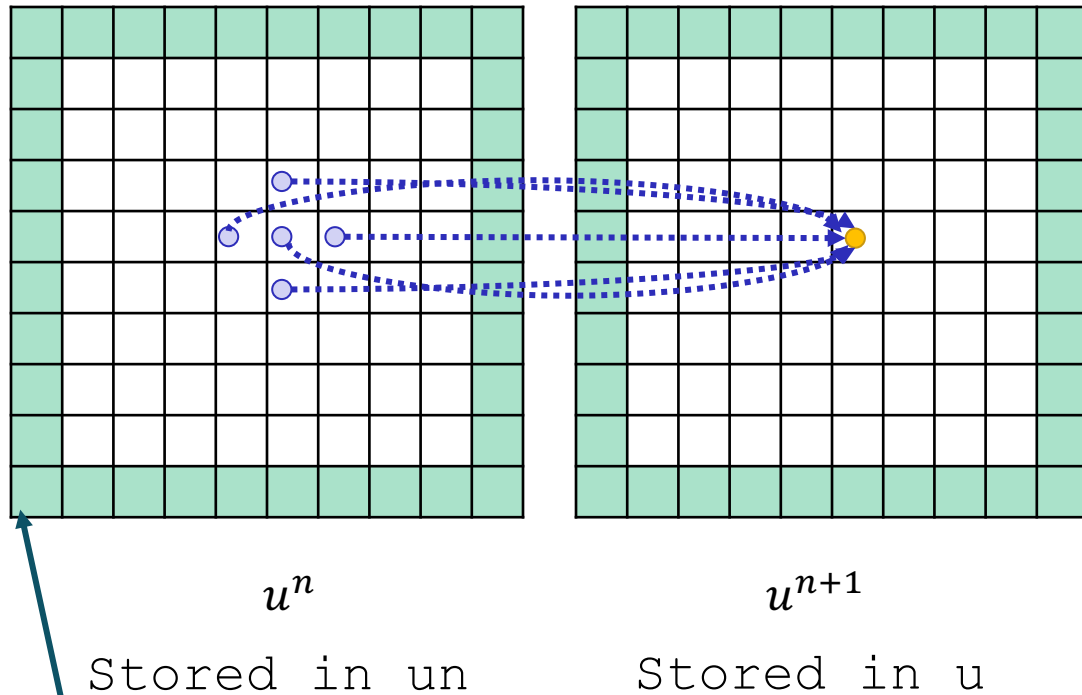
$n = 0$



$n = N$

```
1  const omega = {0.. $\text{nx}$ , 0.. $\text{ny}$ },
2      omegaHat = omega.expand(-1);
3  var u: [omega] real = 1.0;
4  u[nx/4.. $\text{nx}$ /4] = 2.0;
5  var un = u;
6  for 1.. $\text{N}$  {
7      un <=> u
8      forall (i, j) in omegaHat do
9          u[i, j] = un[i, j] + alpha * (
10              un[i-1, j] + un[i, j-1] +
11              un[i+1, j] + un[i, j+1] -
12              4 * un[i, j]);
13  }
```

PARALLEL 2D HEAT EQUATION



Fixed
boundary
values

- This computation uses a "5 point stencil"
- Each point in 'u' can be computed in parallel
 - this is accomplished using a 'forall' loop

```
7  ...
8  forall (i, j) in omegaHat do
9      u[i, j] = un[i, j] + alpha * (
10         un[i-1, j] + un[i, j-1] +
11         un[i+1, j] + un[i, j+1] -
12         4 * un[i, j]);
13  ...
```

$$u_{i,j}^{n+1} = u_{i,j}^n + \alpha(u_{i-1,j}^n + u_{i,j-1}^n + u_{i+1,j}^n + u_{i,j+1}^n - 4u_{i,j}^n)$$

BLOCK DISTRIBUTED & PARALLEL 2D HEAT EQUATION

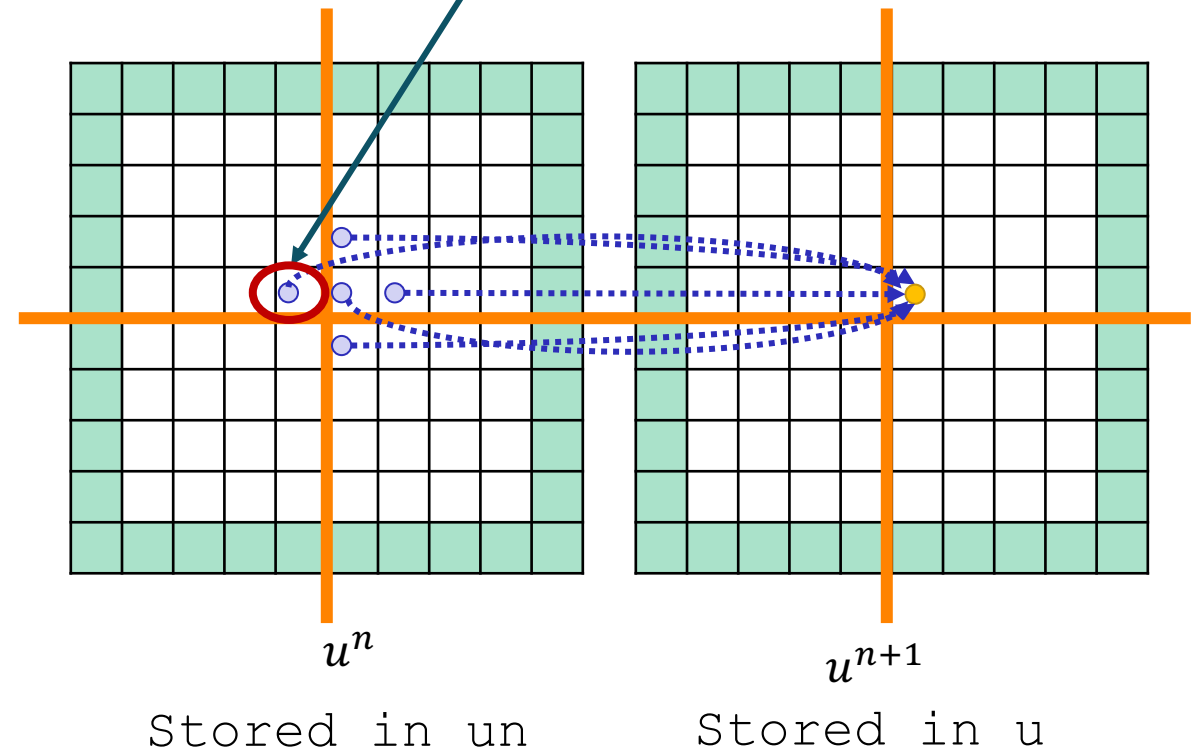
- Declaring distributed domains with the block distribution

```
const Omega = blockDist.createDomain(0.. $\text{nx}$ , 0.. $\text{ny}$ ),  
      OmegaHat = Omega.expand(-1);
```

- Distributed & Parallel loop over 'OmegaHat'

```
for 1.. $\text{nt}$  {  
  u <=> un;  
  
  forall (i, j) in OmegaHat do  
    u[i, j] = un[i, j] + alpha * (  
      un[i-1, j] + un[i, j-1] +  
      un[i+1, j] + un[i, j+1] -  
      4 * un[i, j]);  
}
```

Array access across locale
boundaries automatically
invokes communication



STENCIL DISTRIBUTED & PARALLEL 2D HEAT EQUATION

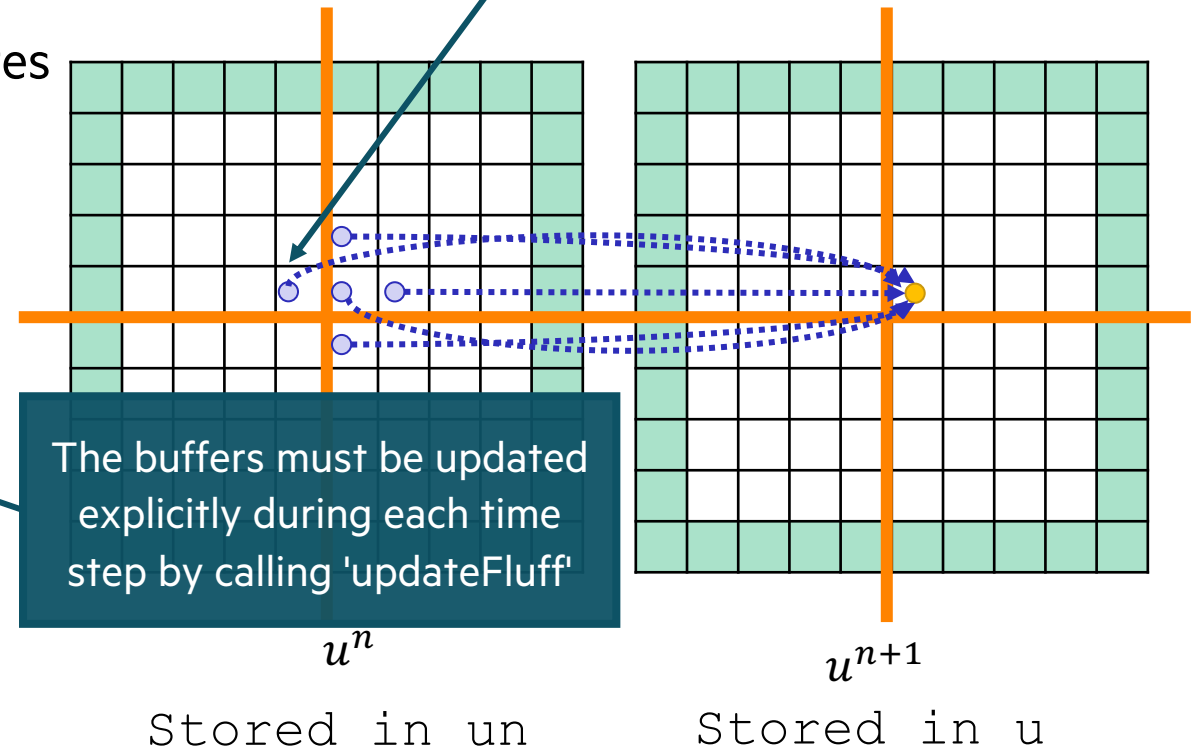
- Declaring distributed domains with the stencil distribution

```
const Omega = stencilDist.createDomain(  
    {0..  
nx, 0..  
ny}, fluff=(1,1)),  
OmegaHat = Omega.expand(-1);
```

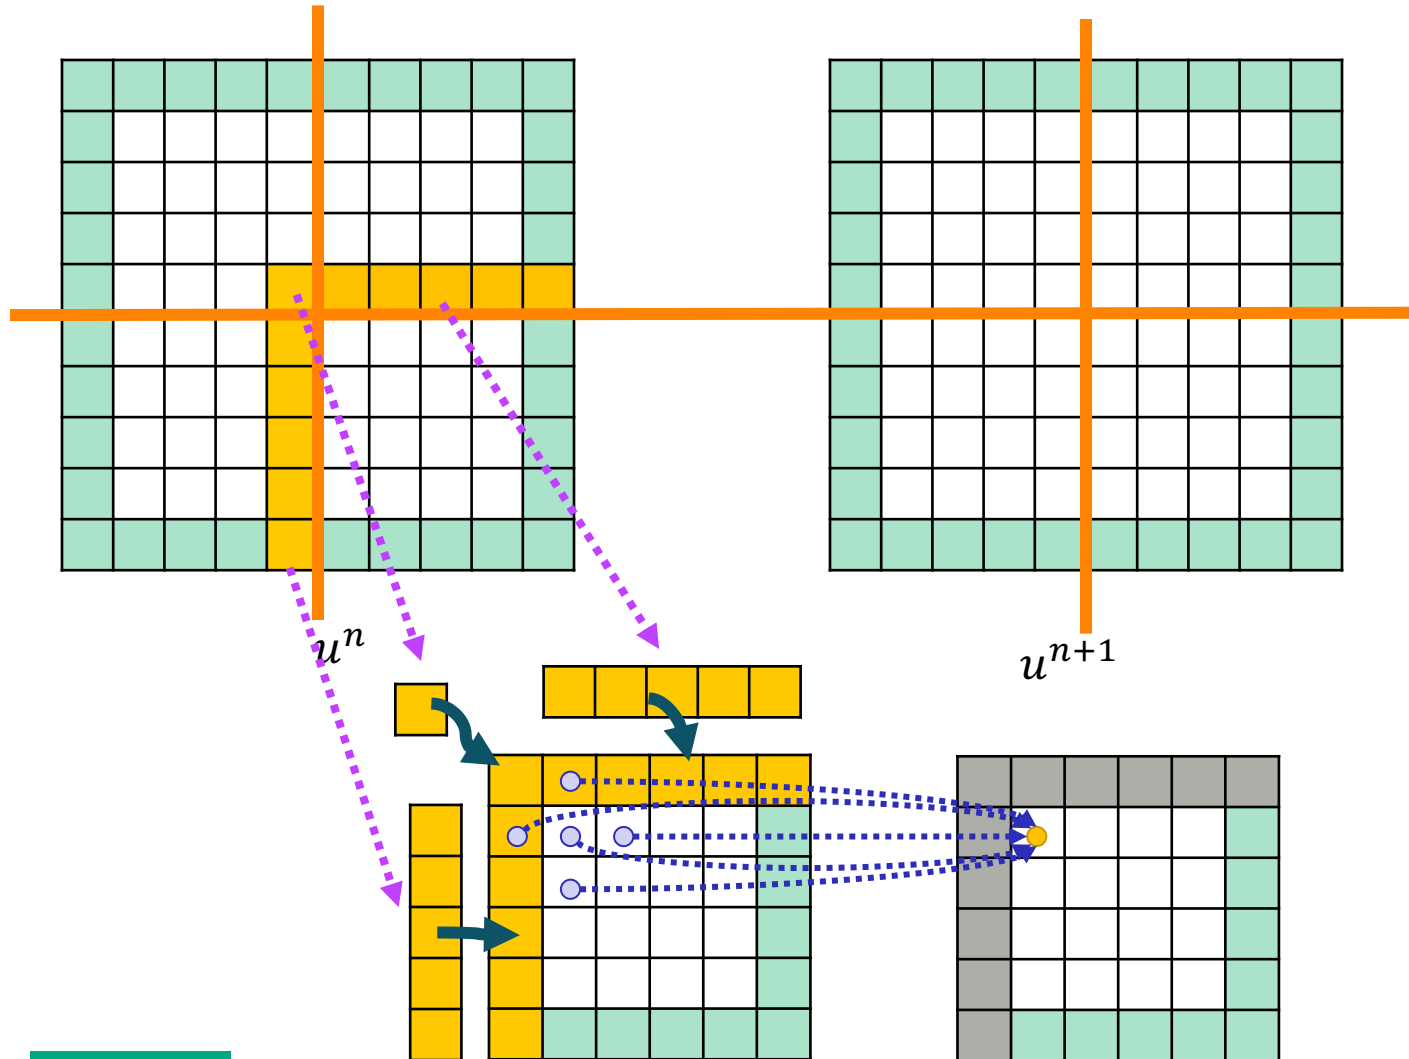
- Distributed & Parallel loop including buffer updates

```
for 1..  
nt {  
    u <=> un;  
    un.updateFluff();  
    forall (i, j) in OmegaHat do  
        u[i, j] = un[i, j] + alpha * (  
            un[i-1, j] + un[i, j-1] +  
            un[i+1, j] + un[i, j+1] -  
            4 * un[i, j]);  
}
```

Array access across locale boundaries (within the fluff region) results in a local buffer access — no communication is required



STENCIL DISTRIBUTED & PARALLEL 2D HEAT EQUATION



- Each locale owns a region of the array surrounded by a "fluff" (buffer) region
- Calling 'updateFluff' copies values from neighboring regions of the array into the local buffered region
- Subsequent accesses of those values result in a local memory access, rather than a remote communication

COMM DIAGNOSTICS

The 'CommDiagnostics' module provides functions for tracking comm between locales

- the following is a common pattern:

```
use CommDiagnostics;  
...  
startCommDiagnostics();  
potentiallyCommHeavyOperation();  
stopCommDiagnostics();  
...  
printCommDiagnosticsTable();
```

- which results in a table summarizing comm counts between the **start** and **stop** calls, e.g.,

locale	get	put	execute_on	execute_on_nb
-----:	--:	--:	-----:	-----:
0	10	0	6	12
1	105	5	0	0
2	105	4	0	0
3	105	7	0	0

- Compiling with '--no-cache-remote' before collecting comm diagnostics is recommended



HANDS ON: HEAT 2D COMM DIAGNOSTICS RESULTS

 heat-2D-block.chpl, heat-2D-stencil.chpl

- Comparing comm diagnostics for:

- heat-2D-block.chpl
- heat-2D-stencil.chpl

- *Compilation:*

```
chpl heat-2D-block.chpl --fast  
      --no-cache-remote -sRunCommDiag=true
```

```
chpl heat-2D-stencil.chpl -fast  
      --no-cache-remote -sRunCommDiag=true
```

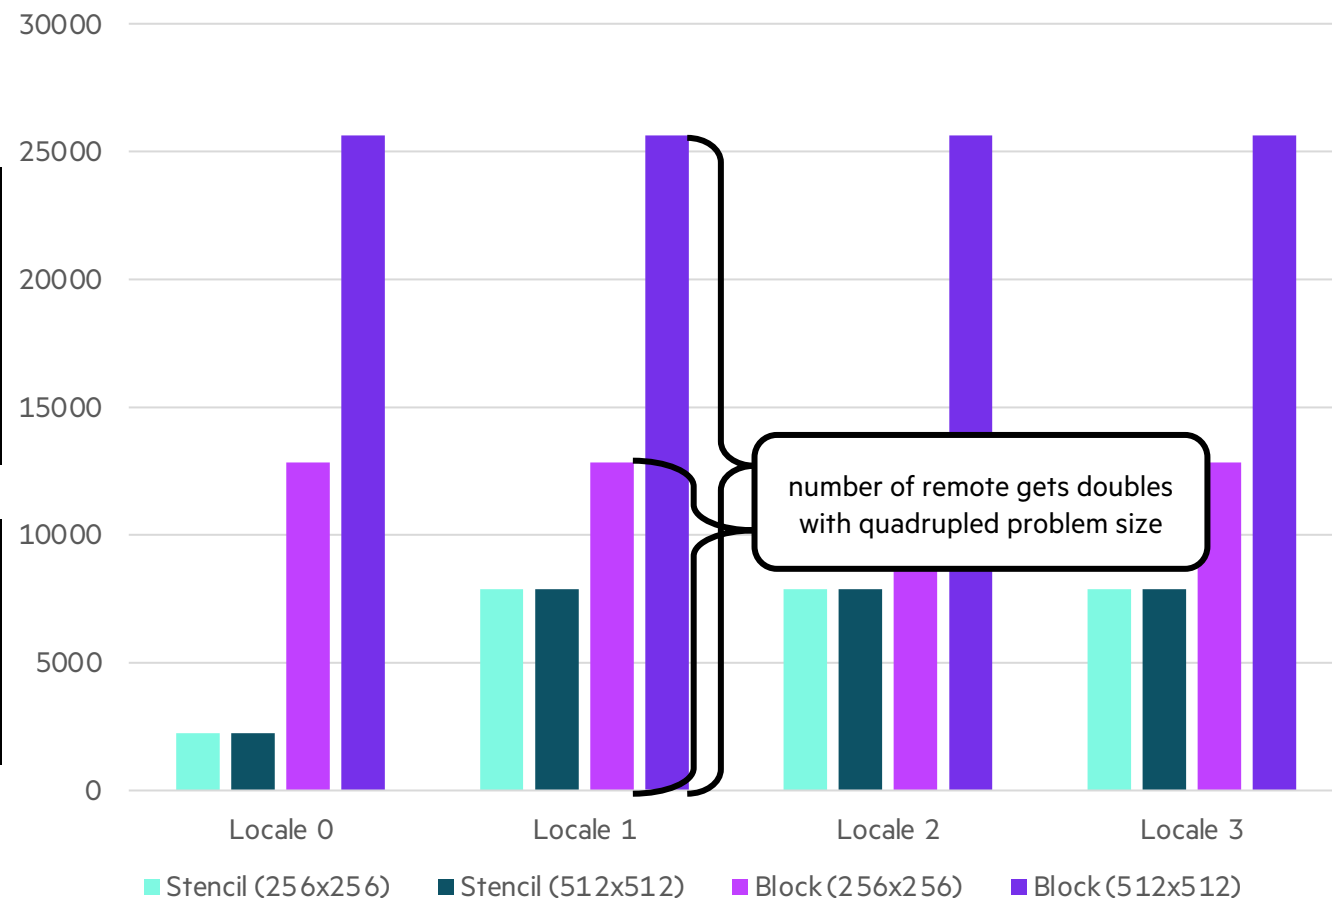
- *Execution:*

```
./heat-2D-block -nl4 --nx=256 --ny=256
```

```
./heat-2D-stencil -nl4 --nx=512 --ny=512
```

- **Block:** number of gets scales with size
- **Stencil:** static number of gets per iteration

Number of Gets on 4 Locales – Block vs. Stencil



OUTLINE: PARALLELISM IN CHAPEL

- Recall processing files in parallel
- Data parallelism concepts and examples including multi-locale parallelism with distributions
- Domains
- Forall Loops
- Domain Distributions
- Using a Different Domain Distribution
- Implicit Communication: Remote writes/Puts and Reads/Gets
- Parallelizing a 1D heat diffusion solver (Hands On)
- Heat 2D example with CommDiagnostics (Hands On)



SUMMARIZING WHAT WE LEARNED IN SESSION 3

- Data parallelism session
 - Provides shared memory and distributed memory parallelism
 - Distributions like block and cyclic can be applied to arrays of any dimension
 - Main control abstraction is the 'forall' loop
 - 'forall' loop uses default iterator over provided array or domain, but can use own iterator
 - This is an example of multi-resolution design in Chapel, i.e., the 'forall' loop is mapped down to lower-level abstractions like 'coforall'
 - CommDiagnostics module can be used to observe the number of remote puts/writes and gets/reads at runtime



LEARNING OBJECTIVES FOR TODAY'S CHAPEL TUTORIAL

- Familiarity with the Chapel execution model including how to run codes in parallel on a single node, across nodes, and both
- Learn Chapel concepts by compiling and running provided code examples
 - ✓ Serial code using map/dictionary, (k-mer counting from bioinformatics)
 - ✓ Parallelism and locality in Chapel
 - ✓ Distributed parallelism and 1D arrays, (processing files in parallel)
 - ✓ Chapel basics in the context of an n-body code
 - ✓ Distributed parallelism and 2D arrays, (heat diffusion problem)
- How to parallelize histogram
- Using CommDiagnostics for counting remote reads and writes
- Chapel and Arkouda best practices including avoiding races and performance gotchas
- Where to get help and how you can participate in the Chapel community



ONE DAY CHAPEL TUTORIAL

- 9-10:30: Getting started using Chapel for parallel programming
- 10:30-10:45: break
- 10:45-12:15: Chapel basics in the context of the n-body example code
- 12:15-1:15: lunch
- 1:15-2:45: Distributed and shared-memory parallelism especially w/arrays (data parallelism)
- 2:45-3:00: break
- 3:00-4:30: More parallelism including for asynchronous parallelism (task parallelism)
- 4:30-5:00: Wrap-up including gathering further questions from attendees



CHAPEL RESOURCES

Chapel homepage: <https://chapel-lang.org>


- (points to all other resources)

Social Media:

- Twitter: [@ChapelLanguage](https://twitter.com/ChapelLanguage)
- Facebook: [@ChapelLanguage](https://facebook.com/ChapelLanguage)
- YouTube: <http://www.youtube.com/c/ChapelParallelProgrammingLanguage>

Community Discussion / Support:

- Discourse: <https://chapel.discourse.group/>
- Gitter: <https://gitter.im/chapel-lang/chapel>
- Stack Overflow: <https://stackoverflow.com/questions/tagged/chapel>
- GitHub Issues: <https://github.com/chapel-lang/chapel/issues>



The Chapel Parallel Programming Language

What is Chapel?

Chapel is a programming language designed for productive parallel computing at scale.

Why Chapel? Because it simplifies parallel programming through elegant support for:

- **distributed arrays** that can leverage thousands of nodes' memories and cores
- **a global namespace** supporting direct access to local or remote variables
- **data parallelism** to trivially use the cores of a laptop, cluster, or supercomputer
- **task parallelism** to create concurrency within a node or across the system

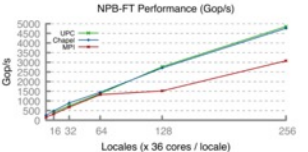
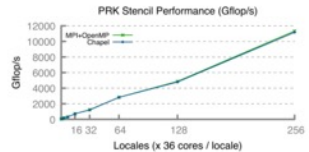
Chapel Characteristics

- **productive:** code tends to be similarly readable/writable as Python
- **scalable:** runs on laptops, clusters, the cloud, and HPC systems
- **fast:** performance competes with or beats C/C++ & MPI & OpenMP
- **portable:** compiles and runs in virtually any *nix environment
- **open-source:** hosted on GitHub, permissively licensed

New to Chapel?

As an introduction to Chapel, you may want to...

- watch an [overview talk](#) or browse its [slides](#)
- read a [blog-length](#) or [chapter-length](#) introduction to Chapel
- learn about [projects powered by Chapel](#)
- check out [performance highlights](#) like these:



PRK Stencil Performance (Gflop/s)

NPB-FT Performance (Gop/s)

- browse [sample programs](#) or [learn](#) how to write distributed programs like this one:

```
use CyclicDist;           // use the Cyclic distribution library
config const n = 100;     // use --n=<val> when executing to override this default

forall i in {1..n} dmapped Cyclic(startIdx=1) do
  writeln("Hello from iteration ", i, " of ", n, " running on node ", here.id);
```