



Chapel: Background



Chapel's Origins

- **HPCS:** High Productivity Computing Systems  
 - Overall goal: Raise high-end user productivity by 10x
Productivity = Performance + Programmability + Portability + Robustness
- **Phase II:** Cray, IBM, Sun (July 2003 – June 2006)
 - Goal: Propose new productive system architectures
 - Each vendor created a new programming language
 - **Cray:** Chapel
 - **IBM:** X10
 - **Sun:** Fortress
- **Phase III:** Cray, IBM (July 2006 –)
 - Goal: Develop the systems proposed in phase II
 - Each vendor implemented a compiler for their language
 - Sun also continued their Fortress effort without HPCS funding

Chapel's Productivity Goals

- Vastly improve **programmability** over current languages
 - Writing parallel programs
 - Reading, modifying, porting, tuning, maintaining them
- Support **performance** at least as good as MPI
 - Competitive with MPI on generic clusters
 - Better than MPI on more capable architectures
- Improve **portability** over current languages
 - As ubiquitous as MPI but more abstract
 - More portable than OpenMP, UPC, and CAF are thought to be
- Improve **robustness** via improved semantics
 - Eliminate common error cases
 - Provide better abstractions to help avoid other errors

Outline

- Chapel's Context
- Chapel's Motivating Themes
 1. General parallel programming
 2. *Global-view* abstractions
 3. *Multiresolution* design
 4. Control over locality/affinity
 5. Reduce gap between mainstream & HPC languages

1) General Parallel Programming

With a unified set of concepts...

...express any parallelism desired in a user's program


- **Styles:** data-parallel, task-parallel, concurrency, nested, ...
- **Levels:** model, function, loop, statement, expression

...target all parallelism available in the hardware

- **Systems:** multicore desktops, clusters, HPC systems, ...
- **Levels:** machines, nodes, cores, instructions


2) Global-View Abstractions

In pictures: “Apply a 3-Point Stencil to a vector”



2) Global-View Abstractions


In pictures: “Apply a 3-Point Stencil to a vector”



2) Global-View Abstractions

In code: “Apply a 3-Point Stencil to a vector”


Global-View



```
def main() {
  var n = 1000;
  var A, B: [1..n] real;

  forall i in 2..n-1 do
    B[i] = (A[i-1] + A[i+1])/2;
  }
```

Local-View (SPMD)



```
def main() {
  var n = 1000;
  var p = numProcs(),
      me = myProc(),
      myN = n/p,
  var A, B: [0..myN+1] real;

  if (me < p-1) {
    send(me+1, A[myN]);
    recv(me+1, A[myN+1]);
  }
  if (me > 0) {
    send(me-1, A[1]);
    recv(me-1, A[0]);
  }


  forall i in 1..myN do
    B[i] = (A[i-1] + A[i+1])/2;
```

Bug: Refers to uninitialized values at ends of A

2) Global-View Abstractions

In code: “Apply a 3-Point Stencil to a vector”

Global-View



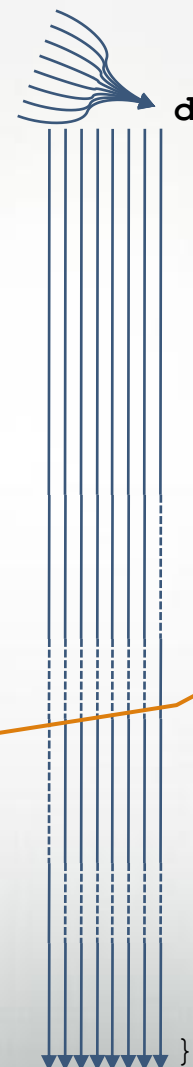
```
def main() {
  var n = 1000;
  var A, B: [1..n] real;

  forall i in 2..n-1 do
    B[i] = (A[i-1] + A[i+1])/2;
  }
```

Communication becomes geometrically more complex for higher-dimensional arrays

Local-View (SPMD)

Assumes p divides n



```
def main() {
  var n = 1000;
  var p = numProcs(),
      me = myProc(),
      myN = n/p,
      iLo = 1,
      iHi = myN;


  var A, B: [0..myN+1] real;

  if (me < p-1) {
    send(me+1, A[myN]);
    rcv(me+1, A[myN+1]);
  } else
    myHi = myN-1;

  if (me > 0) {
    send(me-1, A[1]);
    rcv(me-1, A[0]);
  } else
    myLo = 2;

  forall i in iLo..iHi do
    B[i] = (A[i-1] + A[i+1])/2;
  }
```

2) *rprj3* Stencil from NAS MG





2) *rprj3* Stencil from NAS MG in Chapel

```
def rprj3(S: [?SD], R: [?RD]) {
  const Stencil = [-1..1, -1..1, -1..1],
    W: [0..3] real = (0.5, 0.25, 0.125, 0.0625),
    W3D = [(i,j,k) in Stencil] W[(i!=0) + (j!=0) + (k!=0)];

  forall ijk in SD do
    S[ijk] = + reduce [offset in Stencil]
                  (W3D[offset] * R[ijk + RD.stride*offset]);
}
```



Our previous work in ZPL demonstrated that such compact codes can result in better performance than Fortran + MPI while also supporting more flexibility at runtime*.

*e.g., the Fortran + MPI *rprj3* code shown previously not only assumes p divides n , it also assumes that p and n are specified at compile-time and powers of two.

2) Classifying Current Programming Models

	System	Data Model	Control Model
Communication Libraries	MPI/MPI-2	Local-View	Local-View
	SHMEM, ARMCI, GASNet	Local-View	SPMD
Shared Memory	OpenMP, Pthreads	Global-View (trivially)	Global-View (trivially)
PGAS Languages	Co-Array Fortran	Local-View	SPMD
	UPC	Global-View	SPMD
	Titanium	Local-View	SPMD
PGAS Libraries	Global Arrays	Global-View	SPMD

2) Classifying Current Programming Models

	System	Data Model	Control Model
Communication Libraries	MPI/MPI-2	Local-View	Local-View
	SHMEM, ARMCI, GASNet	Local-View	SPMD
Shared Memory	OpenMP, Pthreads	Global-View (trivially)	Global-View (trivially)
PGAS Languages	Co-Array Fortran	Local-View	SPMD
	UPC	Global-View	SPMD
	Titanium	Local-View	SPMD
PGAS Libraries	Global Arrays	Global-View	SPMD
HPCS Languages	Chapel	Global-View	Global-View
	X10 (IBM)	Global-View	Global-View
	Fortress (Sun)	Global-View	Global-View

2) Global-View Programming: A Final Note

- A language may support both global- and local-view programming — in particular, Chapel does


```

def main() {
    coforall loc in Locales do
        on loc do
            MySPMDProgram(loc.id, Locales.numElements);
}

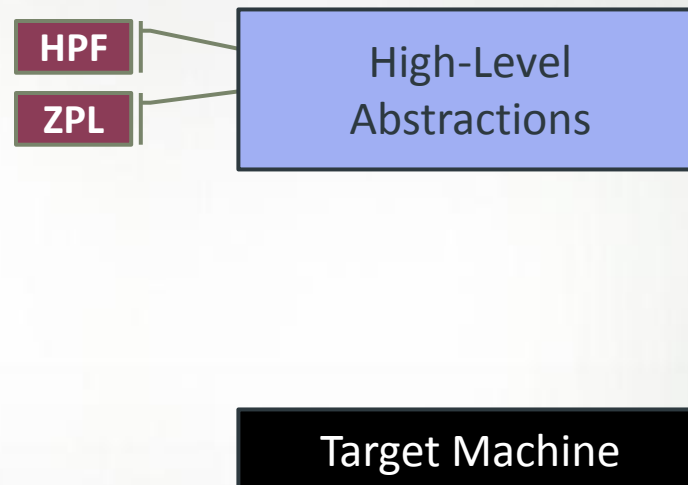
def MySPMDProgram(me, p) {
    . . .
}

```

3) Multiresolution Language Design: Motivation



"Why is everything so difficult?"
"Why don't my programs port trivially?"




"Why don't I have more control?"

3) Multiresolution Language Design

Multiresolution Design: Support multiple tiers of features

- higher levels for programmability, productivity
- lower levels for performance, control
- build the higher-level concepts in terms of the lower

Chapel language concepts



- separate concerns appropriately for clean design

4) Control over Locality/Affinity

Consider:

- Scalable systems tend to store memory with processors
- Remote accesses take longer than local accesses

Therefore:

- Placement of data relative to computation affects scalability
- Programmers need control over data and task placement

Note:

- As core counts grow, locality will matter more on desktops
- GPUs and accelerators already expose node-level locality

5) Reduce Gap Between HPC & Mainstream Languages

Consider:

- Students graduate with training in Java, Matlab, Perl, Python
- Yet HPC programming is dominated by Fortran, C/C++, MPI

We'd like to narrow this gulf with Chapel:

- to leverage advances in modern language design
- to better utilize the skills of the entry-level workforce...
- ...while not ostracizing the traditional HPC programmer
 - e.g., support object-oriented programming, but make it optional

Other examples:

- function overloading, name-based argument passing
- scripting-like features: type inference, generic functions
- rich data structures with iterators (*e.g.*, associative arrays)

Questions?

- Chapel's Context
- Chapel's Motivating Themes
 1. General parallel programming
 2. *Global-view* abstractions
 3. *Multiresolution* design
 4. Control over locality/affinity
 5. Reduce gap between mainstream & HPC languages