

Chapel: Data Parallelism



Data vs. Task Parallelism (Our Terminology)

Data Parallelism:

- parallelism is driven by collections of data
 - data aggregates (arrays)
 - sets of indices (ranges, domains)
 - other user-defined collections
- e.g., “for all elements in array A ...”

Task Parallelism:

- parallelism is expressed using distinct tasks
- e.g., “create a task to do foo() while another does bar()”

(Of course, data parallelism is executed using tasks and task parallelism typically operates on data, so the line can get fuzzy at times...)

"Hello World" in Chapel: a Data Parallel Version

- Data Parallel Hello World

```
config const numIters = 100000;

forall i in 1..numIters do
  writeln("Hello, world! ",
    "from iteration ", i, " of ", numIters);
```

Outline

- Domains and Arrays
 - Regular Domains and Arrays
 - Iterations and Operations
- Other Domain Types
- Reductions and Scans
- NAS MG Stencil Revisited

Domains


Domain: A first-class index set

- Fundamental Chapel concept for data parallelism
- A generalization of ZPL's *region* concept
- Domains may optionally be distributed

Sample Domains

```
config const m = 4, n = 8;

var D: domain(2) = [1..m, 1..n];
```



D


Sample Domains

```

config const m = 4, n = 8;

var D: domain(2) = [1..m, 1..n];

var InnerD: subdomain(D) = [2..m-1, 2..n-1];
  
```



Domains Define Arrays

- Syntax


```
array-type:
  [ domain-expr ] elt-type
```

- Semantics


- Stores element for each index in *domain-expr*

- Example

```
var A, B: [D] real;
```



A



B

- Revisited example

```
var A: [1..3] int; // creates anonymous domain [1..3]
```


Domain Iteration

- For loops (discussed already)
 - Execute loop body once per domain index, serially

```
for i in InnerD do ...
```

	1	2	3	4	5	6	
	7	8	9	10	11	12	

- Forall loops

- Executes loop body once per domain index, in parallel
- Loop must be *serializable* (executable by one task)

```
forall i in InnerD do ...
```

- Loop variables take on **const** domain index values

Other Forall Loops

Forall loops also support...

- A shorthand notation:

```
[ (i,j) in D ] A(i,j) = i + j/10.0;
```

- Expression-based forms:

```
A = forall (i,j) in D do i + j/10.0;
```

```
A = [ (i,j) in D ] i + j/10.0;
```

1.1	1.2	1.3	1.4	1.5	1.6	1.7	1.8
2.1	2.2	2.3	2.4	2.5	2.6	2.7	2.8
3.1	3.2	3.3	3.4	3.5	3.6	3.7	3.8
4.1	4.2	4.3	4.4	4.5	4.6	4.7	4.8

A

Domain Algebra

Domain values support...

- Methods for creating new domains





```
var D2 = InnerD.expand(1,0);
```

- Overloaded Operators

```
var D3 = InnerD + (0,1);
```

- Intersection via Slicing


```
var D4 = D2[D3];
```



Array Slicing/Sub-Arrays

Indexing into arrays with domain values results in a sub-array expression

```
A[InnerD] = B[InnerD + (0,1)];
```



A


B

Array Reallocation


Reassigning a domain logically reallocates its arrays

- values are preserved for common indices

$D = [1..2*m, 1..2*n];$



A



B

Array Iteration

- Array expressions also support for and forall loops

```
for a in A[InnerD] do ...
```

	1	2	3	4	5	6	
	7	8	9	10	11	12	

```
forall a in A[InnerD] do ...
```

- Array loop variables refer to array values (modifiable)

```
forall (a, (i,j)) in (A, D) do a = i + j/10.0;
```

1.1	1.2	1.3	1.4	1.5	1.6	1.7	1.8
2.1	2.2	2.3	2.4	2.5	2.6	2.7	2.8
3.1	3.2	3.3	3.4	3.5	3.6	3.7	3.8
4.1	4.2	4.3	4.4	4.5	4.6	4.7	4.8

Array Indexing

- Arrays can be indexed using variables of their domain's index type (e.g., tuples) or lists of integers

```
var i = 1, j = 2;  
var ij = (i, j);  
  
A[ij] = 1.0;  
A[i, j] = 2.0;
```

- Array indexing can use either parentheses or brackets

```
A(ij) = 3.0;  
A(i, j) = 4.0;
```

Array Arguments and Aliases

- Array values are passed by reference

```
def zero(X: []) { X = 0; }

zero(A[InnerD]); // zeroes the inner values of A
```

- Formal array arguments can reindex actuals

```
def f(X: [1..b,1..b]) { ... } // X uses 1-based indices

f(A[1o..#b, 1o..#b]);
```

- Array alias declarations provide similar functionality

```
var InnerA => A[InnerD];
var InnerA1: [1..n-2,1..m-2] => A[2..n-1,2..m-1];
```


Promoted Functions and Operators

Functions/operators expecting scalars can also take...

- Arrays, causing each element to be passed in

$$\begin{array}{l} \sin(A) \\ 2 * A \end{array} \approx \begin{array}{l} \text{forall } a \text{ in } A \text{ do } \sin(a) \\ \text{forall } a \text{ in } A \text{ do } 2 * a \end{array}$$

- Domains, causing each index to be passed in

$$\text{foo}(\text{Sparse}) \approx \text{forall } i \text{ in } \text{Sparse} \text{ do } \text{foo}(i)$$

Multiple arguments can promote using either...

- Zipper promotion

$$\text{pow}(A, B) \approx \text{forall } (a, b) \text{ in } (A, B) \text{ do } \text{pow}(a, b)$$

- Tensor product promotion

$$\text{pow}[A, B] \approx \text{forall } (a, b) \text{ in } [A, B] \text{ do } \text{pow}(a, b)$$

Where is the Parallelism?

- forall loops are implemented using multiple tasks
 - details vary depending on what is driving the loop
- so are operations that are equivalent to forall loops
 - promoted operators/functions, whole array assignment, ...
- many times, this parallelism can seem invisible
 - for this reason, Chapel's data parallelism can be considered *implicitly parallel*
 - it also tends to make the data parallel features easier to use and less likely to result in bugs as compared to explicit tasks

How Much Parallelism?

By default*, controlled by three configuration variables:

--dataParTasksPerLocale=#

- Specify # of tasks to execute forall loops
- *Current Default:* number of cores

--dataParIgnoreRunningTasks=[true | false]

- If false, reduce # of forall tasks by # of running tasks
- *Current Default:* true

--dataParMinGranularity=#

- If > 0, reduce # of forall tasks if any task has fewer iterations
- *Current Default:* 1

*Default values can be overridden by domain map arguments


Outline

- Domains and Arrays
- Other Domain Types
 - Strided
 - Sparse
 - Associative
 - Opaque
- Reductions and Scans
- NAS MG Stencil Revisited


Chapel Domain Types

Chapel supports several domain types...


```
var OceanSpace = [0..#lat, 0..#long],
    AirSpace = OceanSpace by (2,4),
    IceSpace: sparse subdomain(OceanSpace) = genCaps();
```



dense




strided




sparse

unstructured



associative







```
var Vertices: domain(opaque) = ...,    People: domain(string) = ...;
```

Chapel Array Types

All domain types can be used to declare arrays...

```
var Ocean: [OceanSpace] real,
    Air: [AirSpace] real,
    IceCaps[IceSpace] real;
```



	"steve"
	"lee"
	"sung"
	"david"
	"jacob"
	"albert"
	"brad"






```
var Weight: [Vertices] real,
```

```
Age: [People] int;
```

Iteration

...to iterate over index sets...

```
forall ij in AirSpace do
  Ocean(ij) += IceCaps(ij);
```








```
forall v in Vertices do
  Weight(v) = numEdges(v);
```

```
forall p in People do
  Age(p) += 1;
```

Slicing

...to slice arrays...

```
Ocean[AirSpace] += IceCaps[AirSpace];
```



...Vertices[Interior]...





...People[Interns]...

Reallocation

...and to reallocate arrays

```
AirSpace = OceanSpace by (2,2);
```

```
IceSpace += genEquator();
```



	"steve"
	"lee"
	"sung"
	"david"
	"jacob"
	"albert"
	"brad"
	"vass"

```
newnode = Vertices.create();
```

```
People += "vass";
```

Associative Domains and Arrays by Example

```

var Presidents: domain(string) =
    ("George", "John", "Thomas",
     "James", "Andrew", "Martin");

Presidents += "William";

var Age: [Presidents] int,
    Birthday: [Presidents] string;

Birthday("George") = "Feb 22";

forall president in President do
    if Birthday(president) == today then
        Age(president) += 1;
  
```

George
John
Thomas
James
Andrew
Martin
William

Presidents

Feb 22
Oct 30
Apr 13
Mar 16
Mar 15
Dec 5
Feb 9

Birthday

277
274
266
251
242
227
236

Age

Outline

- Domains and Arrays
- Other Domain Types
- Reductions and Scans
 - Reductions
 - Scans
- NAS MG Stencil Revisited

Reductions

- Syntax

```
reduce-expr:
  reduce-op reduce iterator-expr
```

- Semantics

- Combines argument values using *reduce-op*
- *Reduce-op* may be built-in or user-defined

- Examples

```
total = + reduce A;
bigDiff = max reduce [i in InnerD] abs (A(i)-B(i));
(minVal, minLoc) = minloc reduce (A, D);
```

Scans

- Syntax

```
scan-expr:
  scan-op scan iterator-expr
```

- Semantics

- Computes parallel prefix over values using *scan-op*
- *Scan-op* may be any *reduce-op*

- Examples

```
var A, B, C: [1..5] int;  
A = 1; // A: 1 1 1 1 1  
B = + scan A; // B: 1 2 3 4 5  
B(3) = -B(3); // B: 1 2 -3 4 5  
C = min scan B; // C: 1 1 -3 -3 -3
```

Reduction and Scan Operators


- Built-in
 - $+$, $*$, $\&\&$, $||$, $\&$, $|$, $^$, \min , \max
 - \minloc , \maxloc
 - Takes a tuple of values and indices
 - Generates a tuple of the \min/\max value and its index
- User-defined
 - Defined via a class that supplies a set of methods
 - Compiler generates code that calls these methods
 - Based on:

S. J. Deitz, D. Callahan, B. L. Chamberlain, and L. Snyder. *Global-view abstractions for user-defined reductions and scans*. In Proceedings of the Eleventh ACM SIGPLAN Symposium on Principles and Practices of Parallel Programming, 2006.

Outline

- Domains and Arrays
- Other Domain Types
- Reductions and Scans
- NAS MG Stencil Revisited


Revisiting the *rprj3* Stencil from NAS MG




NAS MG Stencil in Chapel Revisited

```
def rprj3(S: [?SD], R: [?RD]) {
  const Stencil = [-1..1, -1..1, -1..1],
    W: [0..3] real = (0.5, 0.25, 0.125, 0.0625),
    W3D = [(i,j,k) in Stencil] W[(i!=0) + (j!=0) + (k!=0)];

  forall ijk in SD do
    S[ijk] = + reduce [offset in Stencil]
                  (W3D[offset] * R[ijk + RD.stride*offset]);
}
```



S



R

Data Parallelism: Status

- Most features implemented and working correctly
- Regular and irregular domains/arrays generating parallelism
- Scalar performance lacking for higher-dimensional domain/array operations
- Implementation of unstructured domains/arrays is correct but inefficient

Future Directions

- Gain more experience with unstructured (graph-based) domains and arrays

Questions?

- Domains and Arrays
 - Regular Domains and Arrays
 - Iterations and Operations
- Other Domain Types
 - Strided
 - Sparse
 - Associative
 - Opaque
- Data Parallel Operations
 - Reductions
 - Scans
- NAS MG stencil revisited