

Chapel: Sample Codes

Outline

- **STREAM and RA HPC Challenge Benchmarks**
 - simple, regular 1D computations
 - results from SC '09 competition
- **AMR Computations**
 - hierarchical, regular computation
- **SSCA #2**
 - unstructured graph computation

- ## Two classes of competition:

- Class 1:** “best performance”
 - Class 2:** “most productive”
 - Judged on:** 50% performance 50% elegance
 - Four recommended benchmarks:** STREAM, RA, FFT, HPL
 - Use of library routines:** discouraged

- ## Why you may care:

- provides an alternative to the top-500’s focus on peak performance

- ## Recent Class 2 Winners:

2008: *performance:* IBM (UPC/X10)


productive: Cray (Chapel), IBM (UPC/X10), Mathworks (Matlab)

2009: *performance:* IBM (UPC+X10)

elegance: Cray (Chapel)

Chapel Implementation Characteristics

Code Size Summary (Source Lines of Code)




HPC Challenge: Chapel Entries (2008-2009)

Benchmark	2008	2009	Improvement
Global STREAM	1.73 TB/s (512 nodes)	10.8 TB/s (2048 nodes)	6.2x
EP STREAM	1.59 TB/s (256 nodes)	12.2 TB/s (2048 nodes)	7.7x
Global RA	0.00112 GUPs (64 nodes)	0.122 GUPs (2048 nodes)	109x
Global FFT	single-threaded single-node	multi-threaded multi- node	multi-node parallel
Global HPL	single-threaded single-node	multi-threaded single- node	single-node parallel


All timings on ORNL Cray XT4:

- 4 cores/node
- 8 GB/node
- no use of library routines

Global STREAM Triad in Chapel (Excerpts)



```
const ProblemSpace: domain(1, int(64))
    dmapped Block([1..m])
    = [1..m];
```



```
var A, B, C: [ProblemSpace] real;
```



```
forall (a,b,c) in (A,B,C) do
    a = b + alpha * c;
```

This loop should eventually be written:
 $A = B + \alpha * C;$
 (and can be today, but performance is worse)

EP STREAM Triad in Chapel (Excerpts)


```
coforall loc in Locales do
```

```
  on loc {
```



```
    local {
```

```
      var A, B, C: [1..m] real;
```



```
    forall (a,b,c) in (A,B,C) do
```


```
      a = b + alpha * c;
```

```
    }
```

```
  }
```

STREAM Triad Performance

Performance of HPCC STREAM Triad (Cray XT4)



Global Random Access in Chapel (Excerpts)

```

const TableDist = new dmap(new Block([0..m-1])),
    UpdateDist = new dmap(new Block([0..N_U-1]));

const TableSpace: domain ... dmapped TableDist = ...,
    Updates: domain ... dmapped UpdateDist = ...;

var T: [TableSpace] uint(64);

forall ( ,r) in (Updates,RAStream()) do
    on TableDist.idxToLocale(r & indexMask) {
        const myR = r;
        local T(myR & indexMask) ^= myR;
    }
  
```

This body should eventually simply be written:


```

on T(r&indexMask) do
    T(r&indexMask) ^= r;
  
```

(and again, can be today, but performance is worse)


Random Access Performance

Performance of HPCC Random Access (Cray XT4)




Random Access Efficiency on 32+ Nodes

Efficiency of HPCC Random Access on 32+ Locales (Cray XT4)

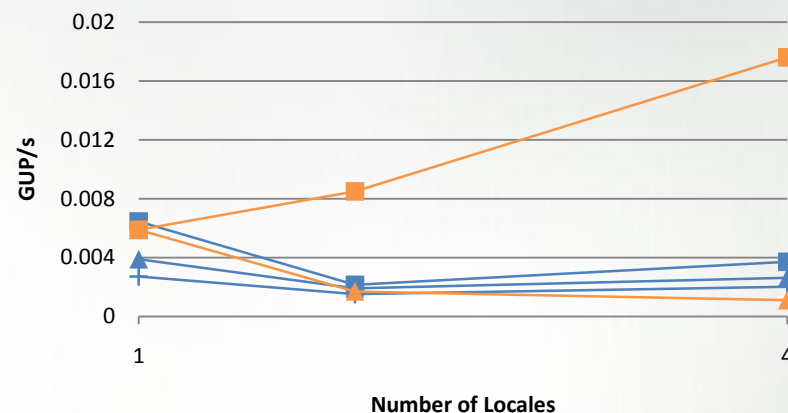


Portability Results


Performance of HPCC STREAM Triad (Cray CX1)



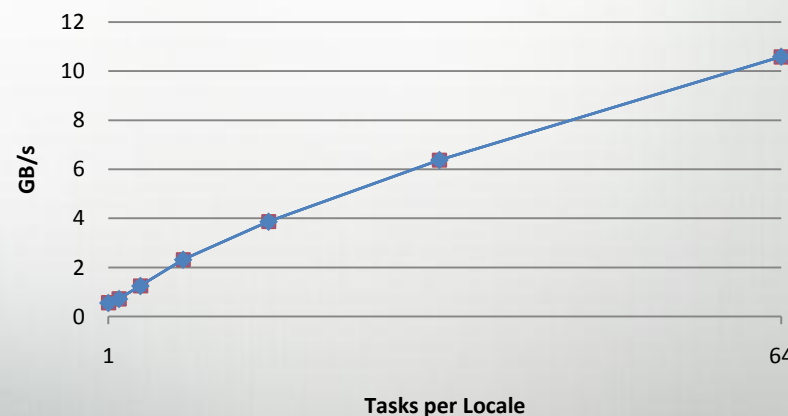
Performance of HPCC Random Access (Cray CX1)



Performance of HPCC STREAM Triad (IBM pSeries 575)



Performance of HPCC STREAM Triad (SGI Altix)




For more on the HPCC Codes

- See `$CHPL_HOME/examples/hpcc` in the release
- Attend the HPC Challenge BoF on Tuesday (12:15-1:15)


Outline

- **STREAM and RA HPC Challenge Benchmarks**
 - simple, regular 1D computations
 - results from SC '09 competition
- **AMR Computations**
 - hierarchical, regular computation
- **SSCA #2**
 - unstructured graph computation

AMR terminology



AMR terminology




Roughly:

Operations on grids




Operations on rectangular
(Chapel) domains


AMR terminology



AMR terminology




AMR terminology



slides courtesy of Jonathan Claridge, UW AMath

AMR terminology



slides courtesy of Jonathan Claridge, UW AMath

Grids: Indexing

- Conventional indexing – number grid cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

Grids: Indexing

- Conventional indexing – number grid cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```

Grids: Indexing

- Conventional indexing – number grid cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```

Rectangular domain: Multidimensional index space

- Supports storage:
var my_array: [cells] real;
- Supports (parallel) iteration:
for(all) cell **in** cells **do** ...

Grids: Indexing

- Conventional indexing – number grid cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```

- Problem with conventional indexing: How are the interfaces indexed?
 - Usual approach: Interface has the same index as the cell above it


Grids: Indexing

- Conventional indexing – number grid cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```

- Problem with conventional indexing: How are the interfaces indexed?
 - Usual approach: Interface has the same index as the cell above it




Grids: Indexing

- Conventional indexing – number grid cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```

- Problem with conventional indexing: How are the interfaces indexed?
 - Usual approach: Interface has the same index as the cell above it




Grids: Indexing

- Conventional indexing – number grid cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```

- Problem with conventional indexing: How are the interfaces indexed?
 - Usual approach: Interface has the same index as the cell above it




Grids: Indexing

- Conventional indexing – number grid cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```

- Problem with conventional indexing: How are the interfaces indexed?
 - Usual approach: Interface has the same index as the cell above it




Grids: Indexing

- Conventional indexing – number grid cells sequentially

3	x	x	x	x
2	x	x	x	x
1	x	x	x	x
	1	2	3	4

```
const cells = [1..4, 1..3];
```


- Problem with conventional indexing: How are the interfaces indexed?
 - Usual approach: Interface has the same index as the cell above it



Many objects will have the same indices


Grids: Indexing

- Modified approach – denser index space



Grids: Indexing


- Modified approach – denser index space



`const cells = [1..7 by 2, 1..5 by 2];`

Grids: Indexing

- Modified approach – denser index space




```
const cells = [1..7 by 2, 1..5 by 2];
```

```
const x_interfaces = [0..8 by 2, 1..5 by 2];
```


Grids: Indexing

- Modified approach – denser index space




```
const cells = [1..7 by 2, 1..5 by 2];
```

```
const x_interfaces = [0..8 by 2, 1..5 by 2];
```

```
const y_interfaces = [1..7 by 2, 0..6 by 2];
```

Grids: Indexing

- Modified approach – denser index space



```
const cells          = [1..7 by 2, 1..5 by 2];  
const x_interfaces = [0..8 by 2, 1..5 by 2];  
const y_interfaces = [1..7 by 2, 0..6 by 2];  
const vertices      = [0..8 by 2, 0..6 by 2];
```


Grids: Indexing

- Modified approach – denser index space



Strided domains



- Array and iteration syntax are **unchanged**
- Chapel helps describe the mathematical problem much more robustly



```
const cells          = [1..7 by 2, 1..5 by 2];  
const x_interfaces = [0..8 by 2, 1..5 by 2];  
const y_interfaces = [1..7 by 2, 0..6 by 2];  
const vertices      = [0..8 by 2, 0..6 by 2];
```

Grids: Dimension independence

- Use the same code to produce results in 2D, 3D, 6D, 17D...



Grids: Dimension independence

- Goal: Use rank-independent domain construction to define a grid of arbitrary spatial dimension

Grids: Dimension independence

- Goal: Use rank-independent domain construction to define a grid of arbitrary spatial dimension
- Begin by setting
`config param dimension: int;`
Specifies a compile-time constant (**param**) that may be specified at the command line (**config**).

Grids: Dimension independence

- Goal: Use rank-independent domain construction to define a grid of arbitrary spatial dimension
- Begin by setting
`config param dimension: int;`
Specifies a compile-time constant (**param**) that may be specified at the command line (**config**).
- A grid is defined by:

```
const x_low, x_high: dimension*real;
```

Coordinate bounds

```
const n_cells: dimension*int;
```

Coordinate bounds

```
const ghost_layer_width: int;
```

Width of ghost cell layer

```
const i_low: dimension*int;
```

Lower index bound

Grids: Dimension independence

- Goal: Use rank-independent domain construction to define a grid of arbitrary spatial dimension
- Begin by setting

```
config param dimension: int;
```

Specifies a compile-time constant (**param**) that may be specified at the command line (**config**).
- A grid is defined by:

```
const x_low, x_high: dimension*real;
```

Coordinate bounds

```
const n_cells: dimension*int;
```

Coordinate bounds

```
const ghost_layer_width: int;
```

Width of ghost cell layer

```
const i_low: dimension*int;
```

Lower index bound

Types `dimension*real` and `dimension*int` are tuples, a native type.

Grids: Dimension independence

- Domain of interior cells:

```
var subranges: dimension*range(stridable=true);
```

```
for d in 1..dimension do
```

```
  subranges(d) = (i_low(d)+1 .. by 2) #n_cells(d);
```

```
var cells:  
cells = cel
```

Unbounded range with correct
lower bound and stride

Count operator: Extracts
n_cells(d) elements

Grids: Dimension independence

- Domain of interior cells:

```
var subranges: dimension*range(stridable=true);
```

```
for d in 1..dimension do
```

```
    subranges(d) = (i_low(d)+1 .. by 2) #n_cells(d);
```

```
var cells: domain(dimension, stridable=true);
```


```
cells = subranges;
```

Define the domain cells

Levels

- Essentially a union of grids


```
var grids: domain(Grid);
```



Levels

- Essentially a union of grids

```
var grids: domain(Grid);
```




Associative domain

- List of indices of *any* type
- Array and iteration syntax are **unchanged**


Levels: Sibling overlaps

- A grid's layer of ghost cells will, in general, overlap some of its siblings. Data will be copied into these overlapped ghost cells prior to mathematical operations.



Levels: Sibling overlaps

- A grid's layer of ghost cells will, in general, overlap some of its siblings. Data will be copied into these overlapped ghost cells prior to mathematical operations.



- Calculating the **overlaps** between siblings:


```
var neighbors: domain(Grid);  
var overlaps: [neighbors] domain(dimension, stridable=true);
```

```
for sibling in neighbors:  
    var overlap = sibling.cells;  
  
    if overlap != this {  
        neighbors.add(sibling);  
        overlaps(sibling) = overlap;  
    }  
}
```

An array of domains; stores one domain for each neighbor.
New space allocated as neighbors grows.

Levels: Sibling overlaps

- A grid's layer of ghost cells will, in general, overlap some of its siblings. Data will be copied into these overlapped ghost cells prior to mathematical operations.



- Calculating the **overlaps** between siblings:


```
var neighbors: domain(Grid);  
var overlaps: [neighbors] domain(dimension, stridable=true);
```

```
for sibling in parent_level.grids {  
    var overlap = 0;  
    for cell in sibling.cells {  
        if overlaps[cell] != this {  
            neighbors[cell] = sibling;  
            overlaps(sibling) = overlap;  
        }  
    }  
}
```

Loop over all grids on the same level, checking for neighbors.

Levels: Sibling overlaps

- A grid's layer of ghost cells will, in general, overlap some of its siblings. Data will be copied into these overlapped ghost cells prior to mathematical operations.



- Calculating the **overlaps** between siblings:

```
var neighbors: domain(Grid);  
var overlaps: [neighbors] domain(dimension, stridable=true);
```

```
for sibling in parent_level.grids {  
  var overlap = extended_cells( sibling.cells );
```


Computes intersection of the domains `extended_cells` and `sibling.cells`.

Take a moment to appreciate what this calculation would look like without domains!

}

Levels: Sibling overlaps

- A grid's layer of ghost cells will, in general, overlap some of its siblings. Data will be copied into these overlapped ghost cells prior to mathematical operations.



- Calculating the **overlaps** between siblings:

```
var neighbors: domain(Grid);  
var overlaps: [neighbors] domain(dimension, stridable=true);
```

```
for sibling in parent_level.grids {  
    var overlap = extended_cells( sibling.cells );
```

```
    if overlap.numIndices > 0 && sibling != this {  
        neighbors.add(sibling);  
        overlaps(sibling) = overlap;  
    }  
}
```


If overlap is nonempty, and sibling is distinct from this grid, then update stored data.

AMR Hierarchy: Brief overview

- Three major challenges
 - Data **coarsening**
 - Data **refinement**
 - **Regridding**


AMR Hierarchy: Brief overview

- Three major challenges
 - Data **coarsening**
 - Data **refinement**
 - **Regridding**
- Coarsening
 - Data transfer occurs on intersection of coarse grid and fine grid




AMR Hierarchy: Brief overview

- Three major challenges
 - Data **coarsening**
 - Data **refinement**
 - **Regridding**
- Coarsening
 - Data transfer occurs on intersection of coarse grid and fine grid




AMR Hierarchy: Brief overview

- Three major challenges
 - Data **coarsening**
 - Data **refinement**
 - **Regridding**
- Coarsening
 - Data transfer occurs on intersection of coarse grid and fine grid




AMR Hierarchy: Brief overview

- Three major challenges
 - Data **coarsening**
 - Data **refinement**
 - **Regridding**
- Coarsening
 - Data transfer occurs on intersection of coarse grid and fine grid
 - Region is rectangular – transfer is relatively easy



AMR Hierarchy: Brief overview

- Three major challenges
 - Data **coarsening**
 - Data **refinement**
 - **Regridding**
- Coarsening
 - Data transfer occurs on intersection of coarse grid and fine grid
 - Region is rectangular – transfer is relatively easy
- Refinement and regridding
 - Involve unions and subtractions of rectangles
 - Much harder; subject of next talk



Conclusion

- Chapel domains make many fundamental AMR calculations very easy, even in a dimension-independent setting

Conclusion

- Chapel domains make many fundamental AMR calculations very easy, even in a dimension-independent setting
- Rectangular domains and associative domains are both very important

Conclusion

- Chapel domains make many fundamental AMR calculations very easy, even in a dimension-independent setting
- Rectangular domains and associative domains are both very important
- Haven't discussed objects for data storage, but Chapel's link between domains and arrays makes them easy to define and use

Conclusion

- Chapel domains make many fundamental AMR calculations very easy, even in a dimension-independent setting
- Rectangular domains and associative domains are both very important
- Haven't discussed objects for data storage, but Chapel's link between domains and arrays makes them easy to define and use

For More on AMR in Chapel


- See <https://chapel.svn.sourceforge.net/svnroot/chapel/trunk/test/studies/amr/>

Outline

- **STREAM and RA HPC Challenge Benchmarks**
 - simple, regular 1D computations
 - results from SC '09 competition
- **AMR Computations**
 - hierarchical, regular computation
- **SSCA #2**
 - unstructured graph computation


SSCA #2 Kernel 2

Given a set of heavy edges *HeavyEdges* in directed graph *G*, find sub-graphs of outgoing paths with $length \leq maxPathLength$



SSCA #2 Kernel 2


Given a set of heavy edges *HeavyEdges* in directed graph G , find sub-graphs of outgoing paths with $length \leq maxPathLength$



maxPathLength = 0

SSCA #2 Kernel 2


Given a set of heavy edges *HeavyEdges* in directed graph G , find sub-graphs of outgoing paths with $length \leq maxPathLength$



$maxPathLength = 0$ $maxPathLength = 1$

SSCA #2 Kernel 2

Given a set of heavy edges *HeavyEdges* in directed graph G , find sub-graphs of outgoing paths with $length \leq maxPathLength$



maxPathLength = 0 *maxPathLength* = 1 *maxPathLength* = 2

SSCA #2 Kernel 2 (Code Courtesy of John Lewis)

```

proc rootedHeavySubgraphs (
    G,
    type vertexSet;
    HeavyEdges      : domain,
    HeavyEdgeSubG    : [],
    in maxPathLength: int ) {
forall (e, subgraph) in
    (HeavyEdges, HeavyEdgeSubG) {
    const (x,y) = e;
    var ActiveLevel: vertexSet;

    ActiveLevel += y;

    subgraph.edges += e;
    subgraph.nodes += x;
    subgraph.nodes += y;
  }
}

```

```

for pathLength in 1..maxPathLength {
    var NextLevel: vertexSet;
    forall v in ActiveLevel do
        forall w in G.Neighbors(v) do
            atomic {
                if !subgraph.nodes.member(w) {
                    NextLevel += w;
                    subgraph.nodes += w;
                    subgraph.edges += (v, w);
                }
            }

    if (pathLength < maxPathLength) then
        ActiveLevel = NextLevel;
  }
}

```

SSCA #2 Kernel 2 (Code Courtesy of John Lewis)

```

proc rootedHeavySubgraphs (
    G,
    type vertexSet;
    HeavyEdges      : domain,
    HeavyEdgeSubG    : [1..
    for pathLength in 1..maxPathLength {
        var NextLevel: vertexSet;
        forall v in ActiveLevel do
            forall w in G.Neighbors(v) do

```

Generic Implementation of Graph G

G.Vertices: A domain whose indices represent the vertices

- For toroidal graphs, a domain(d), so vertices are d -tuples
- For other graphs, a domain(1), so vertices are integers

G.Neighbors: An array over G.Vertices

- For toroidal graphs, a fixed-size array over the domain $[1..2*d]$
- For other graphs...
 - ...an associative domain with indices of type `index(G.vertices)`
 - ...a sparse subdomain of G.Vertices

This kernel and the others are generic w.r.t. these decisions!

```

        nodes.member(w) {
            w;
            es += w;
            es += (v, w);

    PathLength) then
        Level;

```

SSCA #2 Kernel 2 (Code Courtesy of John Lewis)

```
proc rootedHeavySubgraphs (
    G,
    type vertexSet;
```

```
for pathLength in 1..maxPathLength {
    var NextLevel: vertexSet;
    forall v in ActiveLevel do
        forall w in G.Neighbors(v) do
```

Generic with respect to vertex sets

vertexSet: A type argument specifying how to represent vertex subsets

Requirements:

- Parallel iteration
- Ability to add members, test for membership

Options:

- An associative domain over vertices
`domain(index(G.vertices))`
- A sparse subdomain of the vertices
`sparse subdomain(G.vertices)`

```
atomic {
    if !subgraph.nodes.member(w) {
        NextLevel += w;
        subgraph.nodes += w;
        subgraph.edges += (v, w);
    }
}
```

```
if (pathLength < maxPathLength) then
    ActiveLevel = NextLevel;
```

SSCA #2 Kernel 2 (Code Courtesy of John Lewis)

```

proc rootedHeavySubgraphs (
    G,
    type vertexSet;
    HeavyEdges      : domain,
    HeavyEdgeSubG    : [],
    in maxPathLength: int ) {
forall (e, subgraph) in
    (HeavyEdges, HeavyEdgeSubG) {
    const (x,y) = e;
    var ActiveLevel: vertexSet;

```

```

    for pathLength in 1..maxPathLength {
        var NextLevel: vertexSet;
        forall v in ActiveLevel do
            forall w in G.Neighbors(v) do
                atomic {
                    if !subgraph.nodes.member(w) {
                        NextLevel += w;
                        subgraph.nodes += w;
                        subgraph.edges += (v, w);

```

The same genericity applies to subgraphs

```

subgraph.edges += e;
subgraph.nodes += x;
subgraph.nodes += y;

```

```

    if (pathLength < maxPathLength) then
        ActiveLevel = NextLevel;

```

```

    }

```

```

}

```

```

}

```

SSCA #2 Kernel 2 (Code Courtesy of John Lewis)

```

proc rootedHeavySubgraphs (
    G,
    type vertexSet;
    HeavyEdges      : domain,
    HeavyEdgeSubG    : [],
    in maxPathLength: int ) {
forall (e, subgraph) in
    (HeavyEdges, HeavyEdgeSubG) {
    const (x,y) = e;
    var ActiveLevel: vertexSet;

    ActiveLevel += y;

    subgraph.edges += e;
    subgraph.nodes += x;
    subgraph.nodes += y;
  }
}

```

```

for pathLength in 1..maxPathLength {
  var NextLevel: vertexSet;
  forall v in ActiveLevel do
    forall w in G.Neighbors(v) do
      atomic {
        if !subgraph.nodes.member(w) {
          NextLevel += w;
          subgraph.nodes += w;
          subgraph.edges += (v, w);
        }
      }

  if (pathLength < maxPathLength) then
    ActiveLevel = NextLevel;
}
}
}

```

SSCA #2 Kernel 2 (Code Courtesy of John Lewis)

```

proc rootedHeavySubgraphs (
    G,
    type vertexSet;
    HeavyEdges      : domain,
    HeavyEdgeSubG    : [],
    in maxPathLength: int ) {
forall (e, subgraph) in
    (HeavyEdges, HeavyEdgeSubG) {
    const (x,y) = e;
    var ActiveLevel: vertexSet;

    ActiveLevel += y;

    subgraph.edges += e;
    subgraph.nodes += x;
    subgraph.nodes += y;
  }
}

```

```

for pathLength in 1..maxPathLength {
  var NextLevel: vertexSet;
  forall v in ActiveLevel do
    forall w in G.Neighbors(v) do
      atomic {
        if !subgraph.nodes.member(w) {
          NextLevel += w;
          subgraph.nodes += w;
          subgraph.edges += (v, w);
        }
      }

    if (pathLength < maxPathLength) then
      ActiveLevel = NextLevel;
  }
}

```

For more on SSca #2

- See `$CHPL_HOME/examples/ssca2` in release

Questions?

- **STREAM and RA HPC Challenge Benchmarks**
 - simple, regular 1D computations
 - results from SC '09 competition
- **AMR Computations**
 - hierarchical, regular computation
- **SSCA #2**
 - unstructured graph computation